

we usually flatten the input:

from $(64, 1, 28, 28)$ to $(64, 784)$

Build a network for identifying images:

`inputs = images.view(images.shape[0], -1)`

input units

choose the appropriate size

`w1 = torch.randn(784, 256)`

`b1 = torch.randn(256)`

`w2 = torch.randn(256, 10)`

`b2 = torch.randn(10)`

hidden units

one output digit for each digit.

one hidden layer

`h = sigmoid(torch.mm(inputs, w1) + b1)`

`out = torch.mm(h, w2) + b2`

`def softmax(x):`

`return torch.exp(x) / torch.sum(torch.exp(x), dim=1)`

• `view(-1, 1)`

`probs = softmax(out)` $\#$ shape = $(64, 10)$

Neural Networks with Pytorch

from torchvision import datasets, transforms

for
normalizing
data

transform = transforms.Compose([transforms.ToTensor(),
transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

download
mnist
dataset

trainset = datasets.MNIST('MNIST-data/',
download = True,
train = True,
transform = transform)

define a
loader

trainloader = torch.utils.data.DataLoader(trainset,

each time we ← batch-size = 64,
get a batch, it's shuffle = True)

shape is (64, 1, 28, 28)

64
images

1 color
channel

image size

dataiter = iter(trainloader)

images, labels = dataiter.next()