



Training Neural Networks

loss function \rightarrow measure prediction error

$$l = \frac{1}{2n} \sum_i^n (y_i - \hat{y}_i)^2$$

\downarrow true labels \hookrightarrow predicted labels

number of
training examples



Loss in Pytorch:

`nn.CrossEntropyLoss` → `nn.LogSoftmax()`
→ `nn.NLLLoss()`

input: Scores (not probabilities)

In summary

1. Build a feed forward network:

`model = nn.Sequential (nn.Linear(784, 128),`

`nn.ReLU,`

`nn.Linear(128, 64),`

← `nn.ReLU,`

`nn.Linear(64, 10))`

ReLU is almost
always used for
hidden layers

2. Define the loss:

and set this
to `nn.NLLLoss()`

→ we can add a
`nn.LogSoftmax(dim=1)`

`criterion = nn.CrossEntropyLoss()`

3. Get the data:

`images, labels = next(iter(train_loader))`

4. Flatten images:

`images = images.view(images.shape[0], -1)`

5. forward pass — get logits

`logits = model(images)`

6. Calculate the loss:

$$\text{loss} = \text{Criterion}(\text{logits}, \text{labels})$$

Autograd (Backpropagation)

track $\rightarrow x = \text{torch.zeros}(1, \text{requires_grad}=\text{True})$

[with `torch.no_grad()`:
turns off autograd

we can get the grad-fn with `x.grad_fn`
finally, we do `z.backward()`

Loss and Autograd together:

Code as before
:
:
loss =

`loss.backward()` \rightarrow compute gradients

Using gradients: ^{before} forward pass \rightarrow `optimizer.zero_grad()`
we use optimizers : $\left\{ \begin{array}{l} \text{optimizer.step()} \\ \text{after getting the loss} \end{array} \right.$

Summary of Training Steps:

1. forward pass
2. compute loss
3. `loss.backward()` to compute gradients
4. take a step with optimizer to update weights

Full training code:

```
criterion = nn.NLLLoss()  
optimizer = optim.SGD(model.parameters(), lr=0.003)
```

```
for e in range(epochs):
```

```
    running_loss = 0
```

```
    for images, labels in trainloader:
```

```
        images = images.view(images.shape[0], -1)
```

```
        ★ optimizer.zero_grad()
```

```
        Forward pass
```

```
        ★ loss.backward()
```

```
        ★ optimizer.step()
```

```
        running_loss += loss.item()
```

```
    else:
```

```
        training_loss = running_loss / len(trainloader)
```


Inference and Validation

we had this code for testing the prediction:

model = Classifier()

images, labels = next(iter(testloader))

ps = torch.exp(model(images))

└─→ (64, 10)
#images #classes

top_p, top_class = ps.topk(1, dim=1)
 ↓
 k highest value

equals = top_class == labels.view(*top_class.shape)

accuracy = torch.mean>equals.type(torch.FloatTensor))

because
equals is a bit tensor
and we have to convert it.

Validation
step

Implementing Validation Pass:

for e in epochs:

for image, label in trainloader:

as before

else:

with torch.no_grad():

above
code

the
validation
pass