

## Putting Everything Together:

```
device = torch.device("cuda"  
                        if torch.cuda.is_available()  
                        else  
                        "cpu")
```

```
model = models.resnet50(pretrained=True)
```

freeze {

```
for param in model.parameters():  
    param.requires_grad = False
```

replace  
the pre-  
trained  
classifier {

```
classifier = nn.Sequential(nn.Linear(2048, 512),  
                           nn.ReLU(),  
                           nn.Dropout(p=0.2),  
                           nn.Linear(512, 2),  
                           nn.LogSoftmax(dim=1))  
model.fc = classifier
```

```
criterion = nn.NLLLoss()
```

```
optimizer = optim.Adam(model.fc.parameters(),  
                        lr=0.003)
```

move the  
model to ← model.to(device)  
the available  
device.



set hyper params {  
epochs = 1  
steps = 0  
running\_loss = 0  
print\_every = 5

start actual training → for epoch in range(epochs):  
for ims, labs in trainloader:  
steps += 1

move to ← ims, labs = ims.to(device), labs.to(device)  
the available device

optimizer.zero\_grad()

logps = model(images)

loss = criterion(logps, labels)

loss.backward()

optimizer.step()

running\_loss += loss.item()

frequency for validation → if step % print\_every == 0:  
model.eval()  
test\_loss = 0  
accuracy = 0



remember to transfer  
to gpu:

ims, labs = ims.to(device)  
labs.to(device)

Validation

for ims, labs in testloader:

logps = model(ims)

loss = Criterion(logps, labs)

test\_loss += loss.item()

ps = torch.exp(logps)

top\_ps, top\_class = ps.topk(1,  
dim=1)

equality = top\_class ==

labels.view(\*top\_cls.shape)

accuracy += torch.mean(  
equality.type(torch.FloatTensor

)).item()

running\_loss = 0

model.train()