

MNIST problem using PyTorch

```
from torchvision import datasets
```

```
import torchvision.transforms as transforms
```

```
transform = transforms.ToTensor() → Convert data to torch.FloatTensor
```

```
train_data = datasets.MNIST (root = 'data',  
                             train = True,  
                             transform = transform)
```

similarly for test_data (train = False)

```
train_loader = torch.utils.data.DataLoader(  
    train_data,  
    batch_size = b_size,  
    num_workers = workers)
```

Inspecting the data:

one batch of the data {
 dataiter = iter(train_loader)
 images, labels = dataiter.next()
 images = images.numpy()

one image \leftarrow `img = np.squeeze(images[1])`

Defining the Model:

```
import torch.nn as nn
```

```
import torch.nn.functional as F
```

```
class Net(nn.Module):
```

```
    def __init__(self):
```

```
        super(Net, self).__init__()
```

```
        self.fc1 = nn.Linear(28*28, 512)
```

```
        self.fc2 = nn.Linear(512, 512)
```

```
        self.fc3 = nn.Linear(512, 10)
```

to prevent
overfitting \leftarrow

```
        self.dropout = nn.Dropout(0.2)
```

```
    def forward(self, x):
```

```
        # Flatten  $\leftarrow$  x = x.view(-1, 28*28)
```

```
        # apply layer  $\leftarrow$  x = F.relu(self.fc1(x)) ; x = self.dropout(x)  
        # same for fc2.
```

```
        x = self.fc3(x)
```

```
        return x
```


#question: why do we need activation functions?

to scale the outputs of a layer so that they are a consistent, small value.

```
model = Net()
```

```
criterion = nn.CrossEntropyLoss()
```

```
optimizer = optim.Adam(model.parameters(), lr=0.01)
```

Finally, train the model!

```
n_epochs = 30 # start small!
```

```
for epoch in range(n_epochs):
```

```
|     train_loss = 0
```

```
|     for data, target in data_loader:
```

```
|         |     optimizer.zero_grad() → clear all the gradients
```

```
|         |     output = model(data) forward pass
```

```
|         |     loss = criterion(output, target) compute loss
```

```
|         |     loss.backward() → backward propagation
```

```
|         |     optimizer.step() → parameter update
```

```
|         |     train_loss += loss.item() * data.size(0)
```

```
|     train_loss = train_loss / len(train_loader.dataset)
```


test! model.eval()

for data, target in test_loader:

output = model(data)

loss = criterion(output, target)

test_loss += loss.item() * data.size(0)

_, pred = torch.max(output, 1)

correct =

np.squeeze(pred.eq(target.data.view_as(pred)))

★ Question: nn.CrossEntropy, does both of softmax and NLLoss. What if we want the output of the model as class probabilities instead of scores?

Answer: use Softmax and NLLoss separately.

for i in range(batch_size):

label = target.data[i]

class_correct[label] += correct[i].item()

class_total[label] += 1

test_loss = test_loss / len(test_loader.dataset)

for i in range(num_classes):

if class_total[i] > 0:

$$\underline{\text{accuracy}[i]} = \text{np.sum}(\text{class_correct}[i]) / \text{np.sum}(\text{class_total}[i])$$

$$\underline{\text{overall_accuracy}} = \text{np.sum}(\text{class_correct}) / \text{np.sum}(\text{class_total})$$

Model Validation

from torch.utils.data.sampler import SubsetRandomSampler
we use this to sample some validation data

train_sampler = SubsetRandomSampler(train_idx)
valid_sampler = SubsetRandomSampler(valid_idx)

train_loader (as before , sampler=train_sampler)
we do the same for valid_sampler.

Question: why do we need validation?

Answer: knowing when to stop training to avoid overfitting
(when valid-loss starts increasing but train-loss
keep decreasing)

OpenCV

```
import cv2
```

```
import numpy as np
```

```
import matplotlib.image as mpimg
```

load → `image = mpimg.imread('img.png')`

make black/white → `gray = cv2.cvtColor(image, cv2.COLOR_RGB2GRAY)`

the filter → `filt = np.array([[-1, -2, -1],
[0, 0, 0],
[1, 2, 1]])`

convolute → `filtered_image = cv2.filter2D(gray, -1, filt)`