



Components of CNNs

1. number of filters
2. size of each filter
3. stride \rightarrow when sliding, we choose how much we slide at each step.

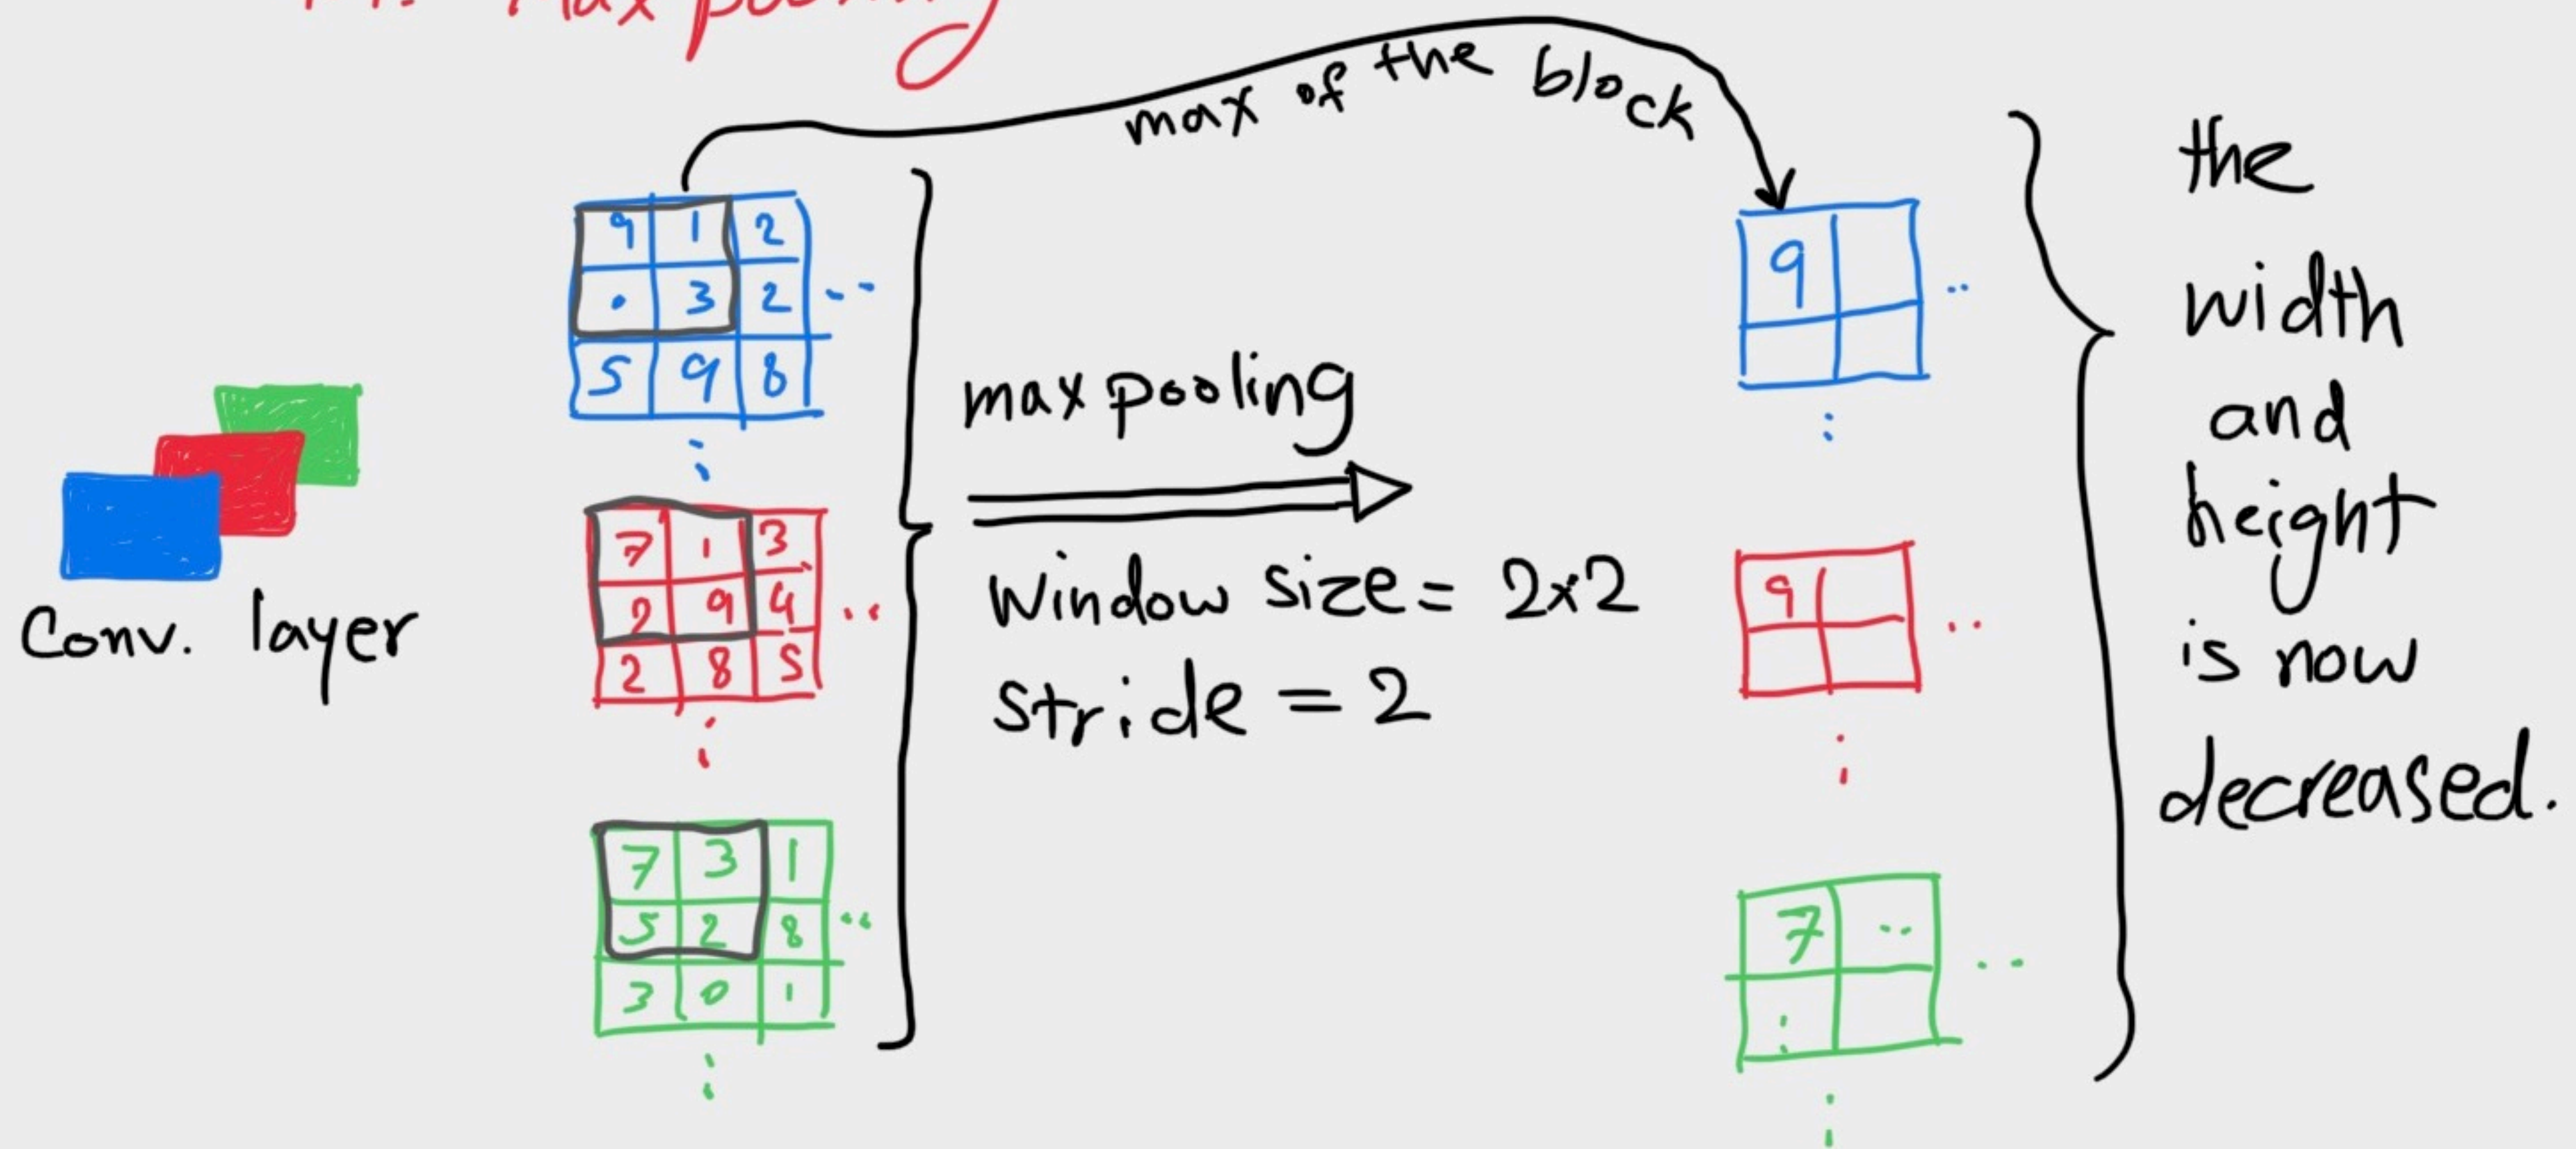
when $\text{stride} = 1$, the filtered image will have the same size as the input image. The size of the filtered image decreases when we increase the stride.

4. pooling layers



Convolutional layers can get too large, i.e., high-param. we use pooling to deal with that.

4.1. Max pooling



4.2. Average pooling

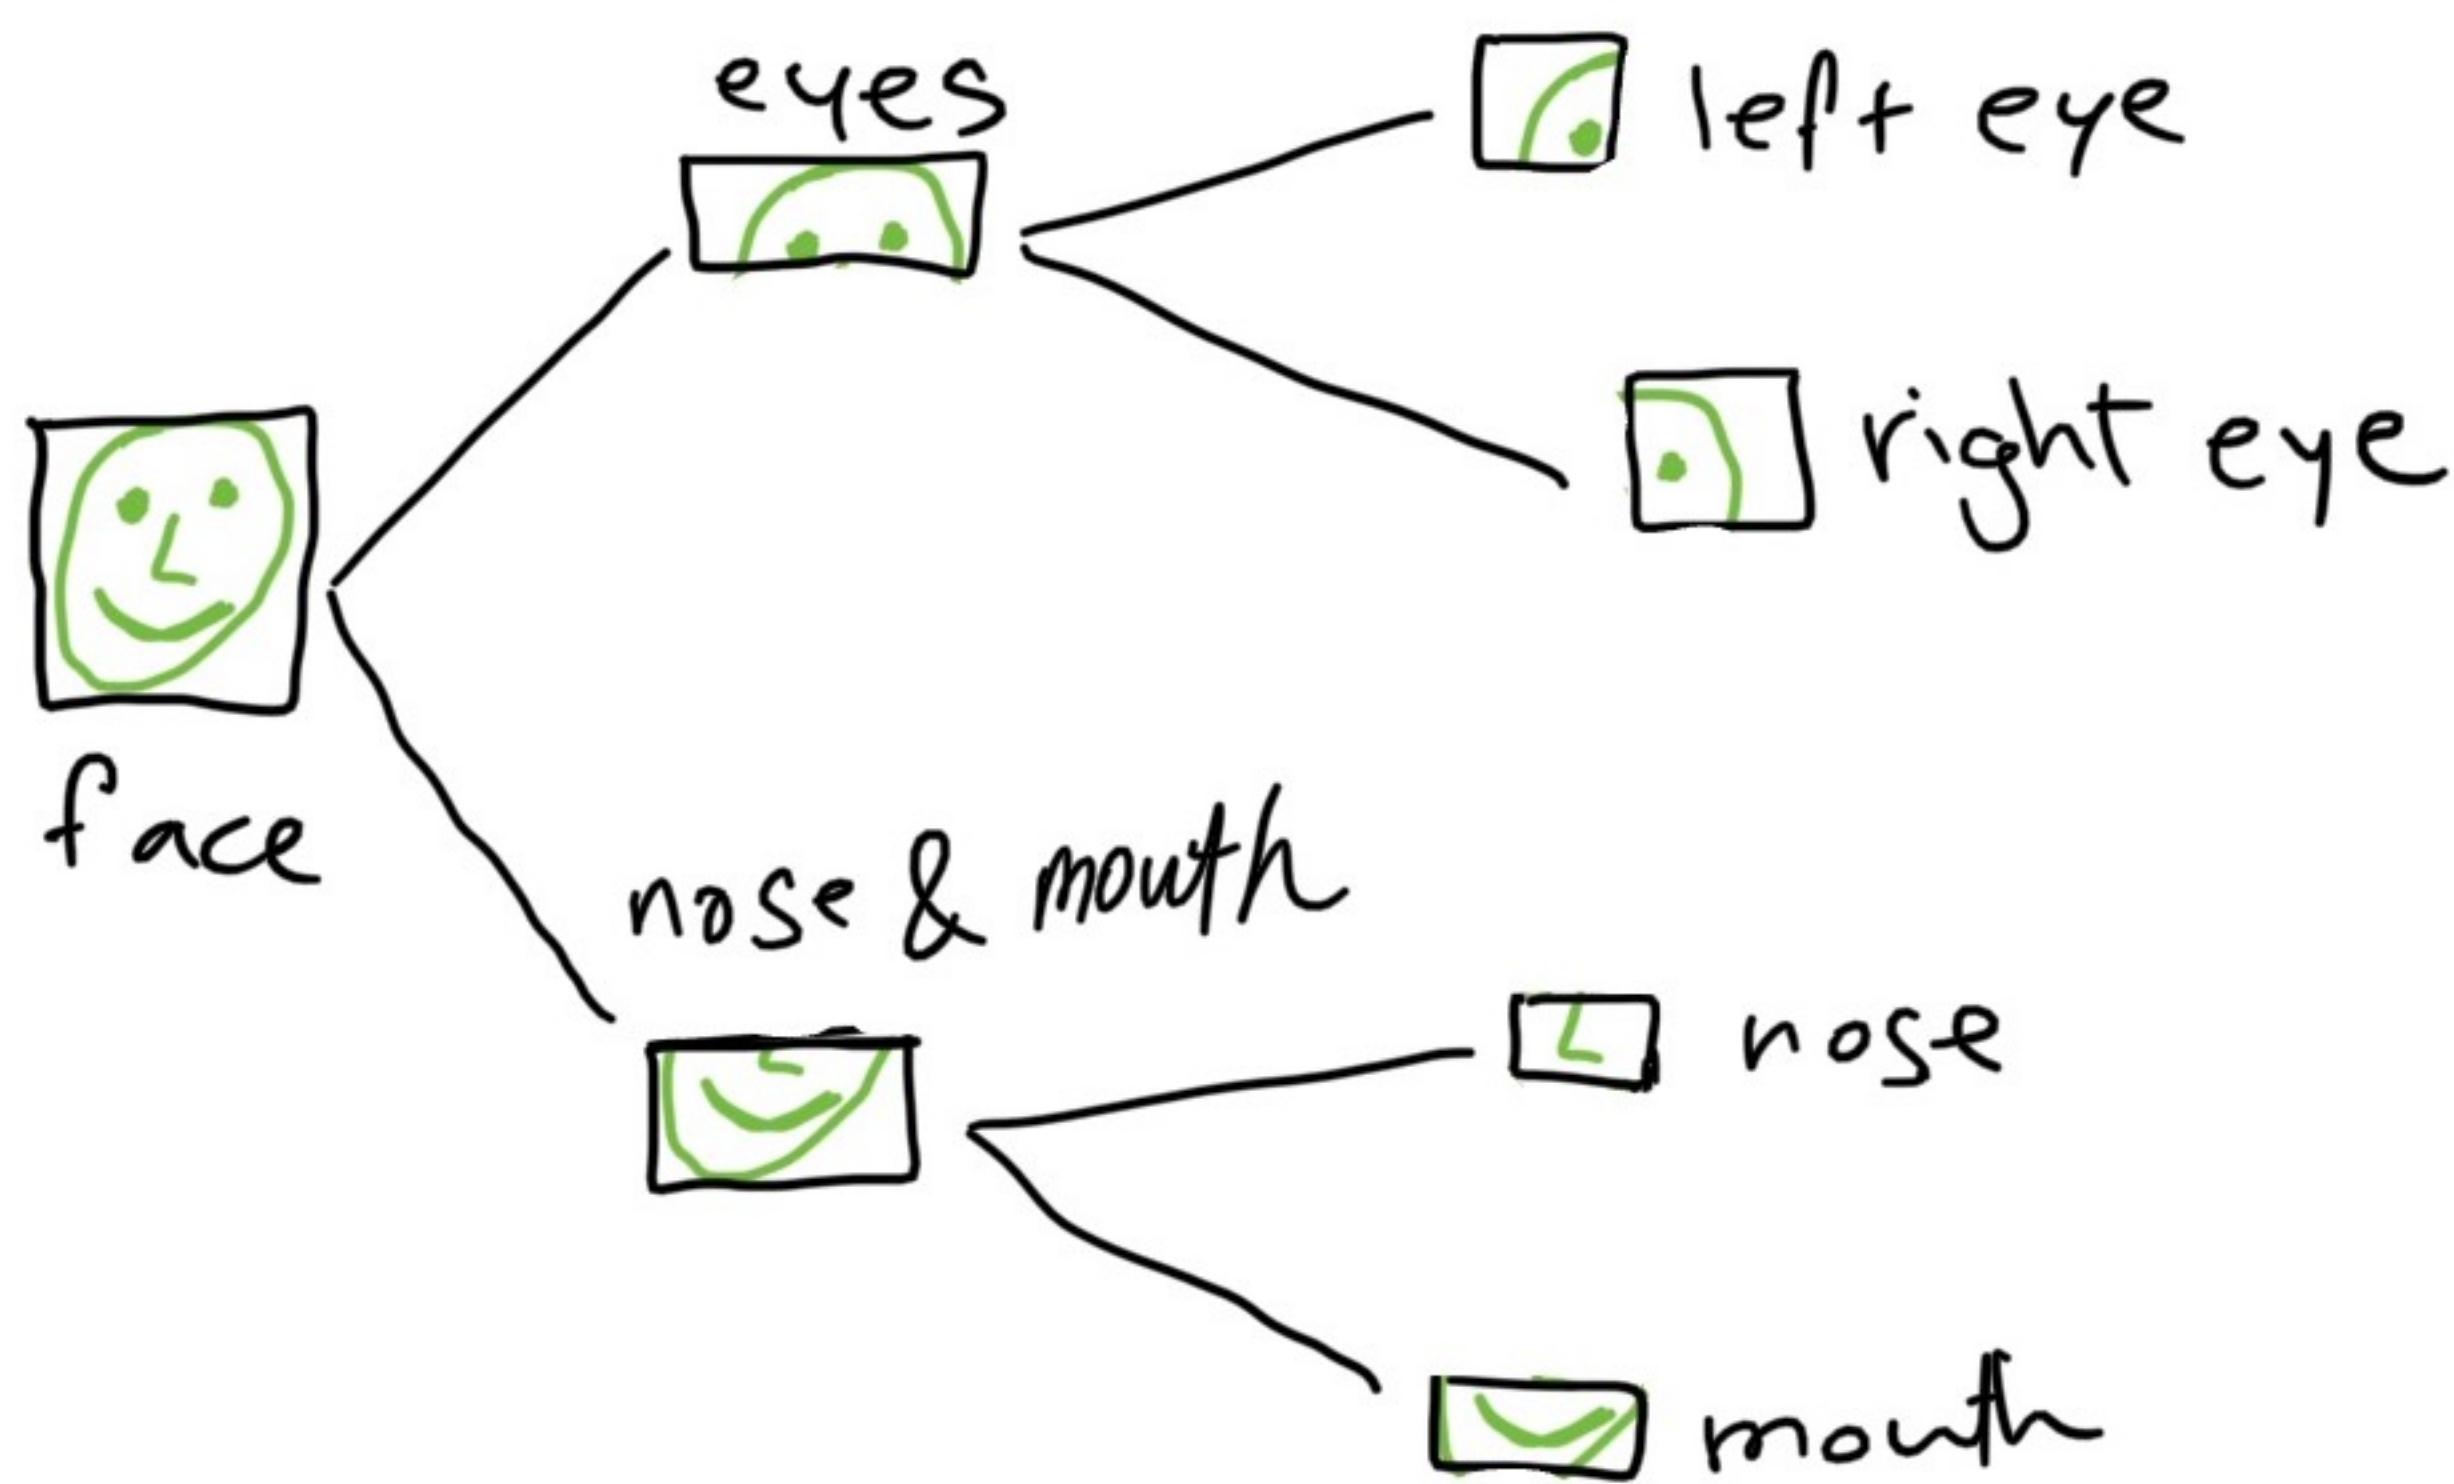
not usually for edge detection

4.3. Alternatives to pooling :

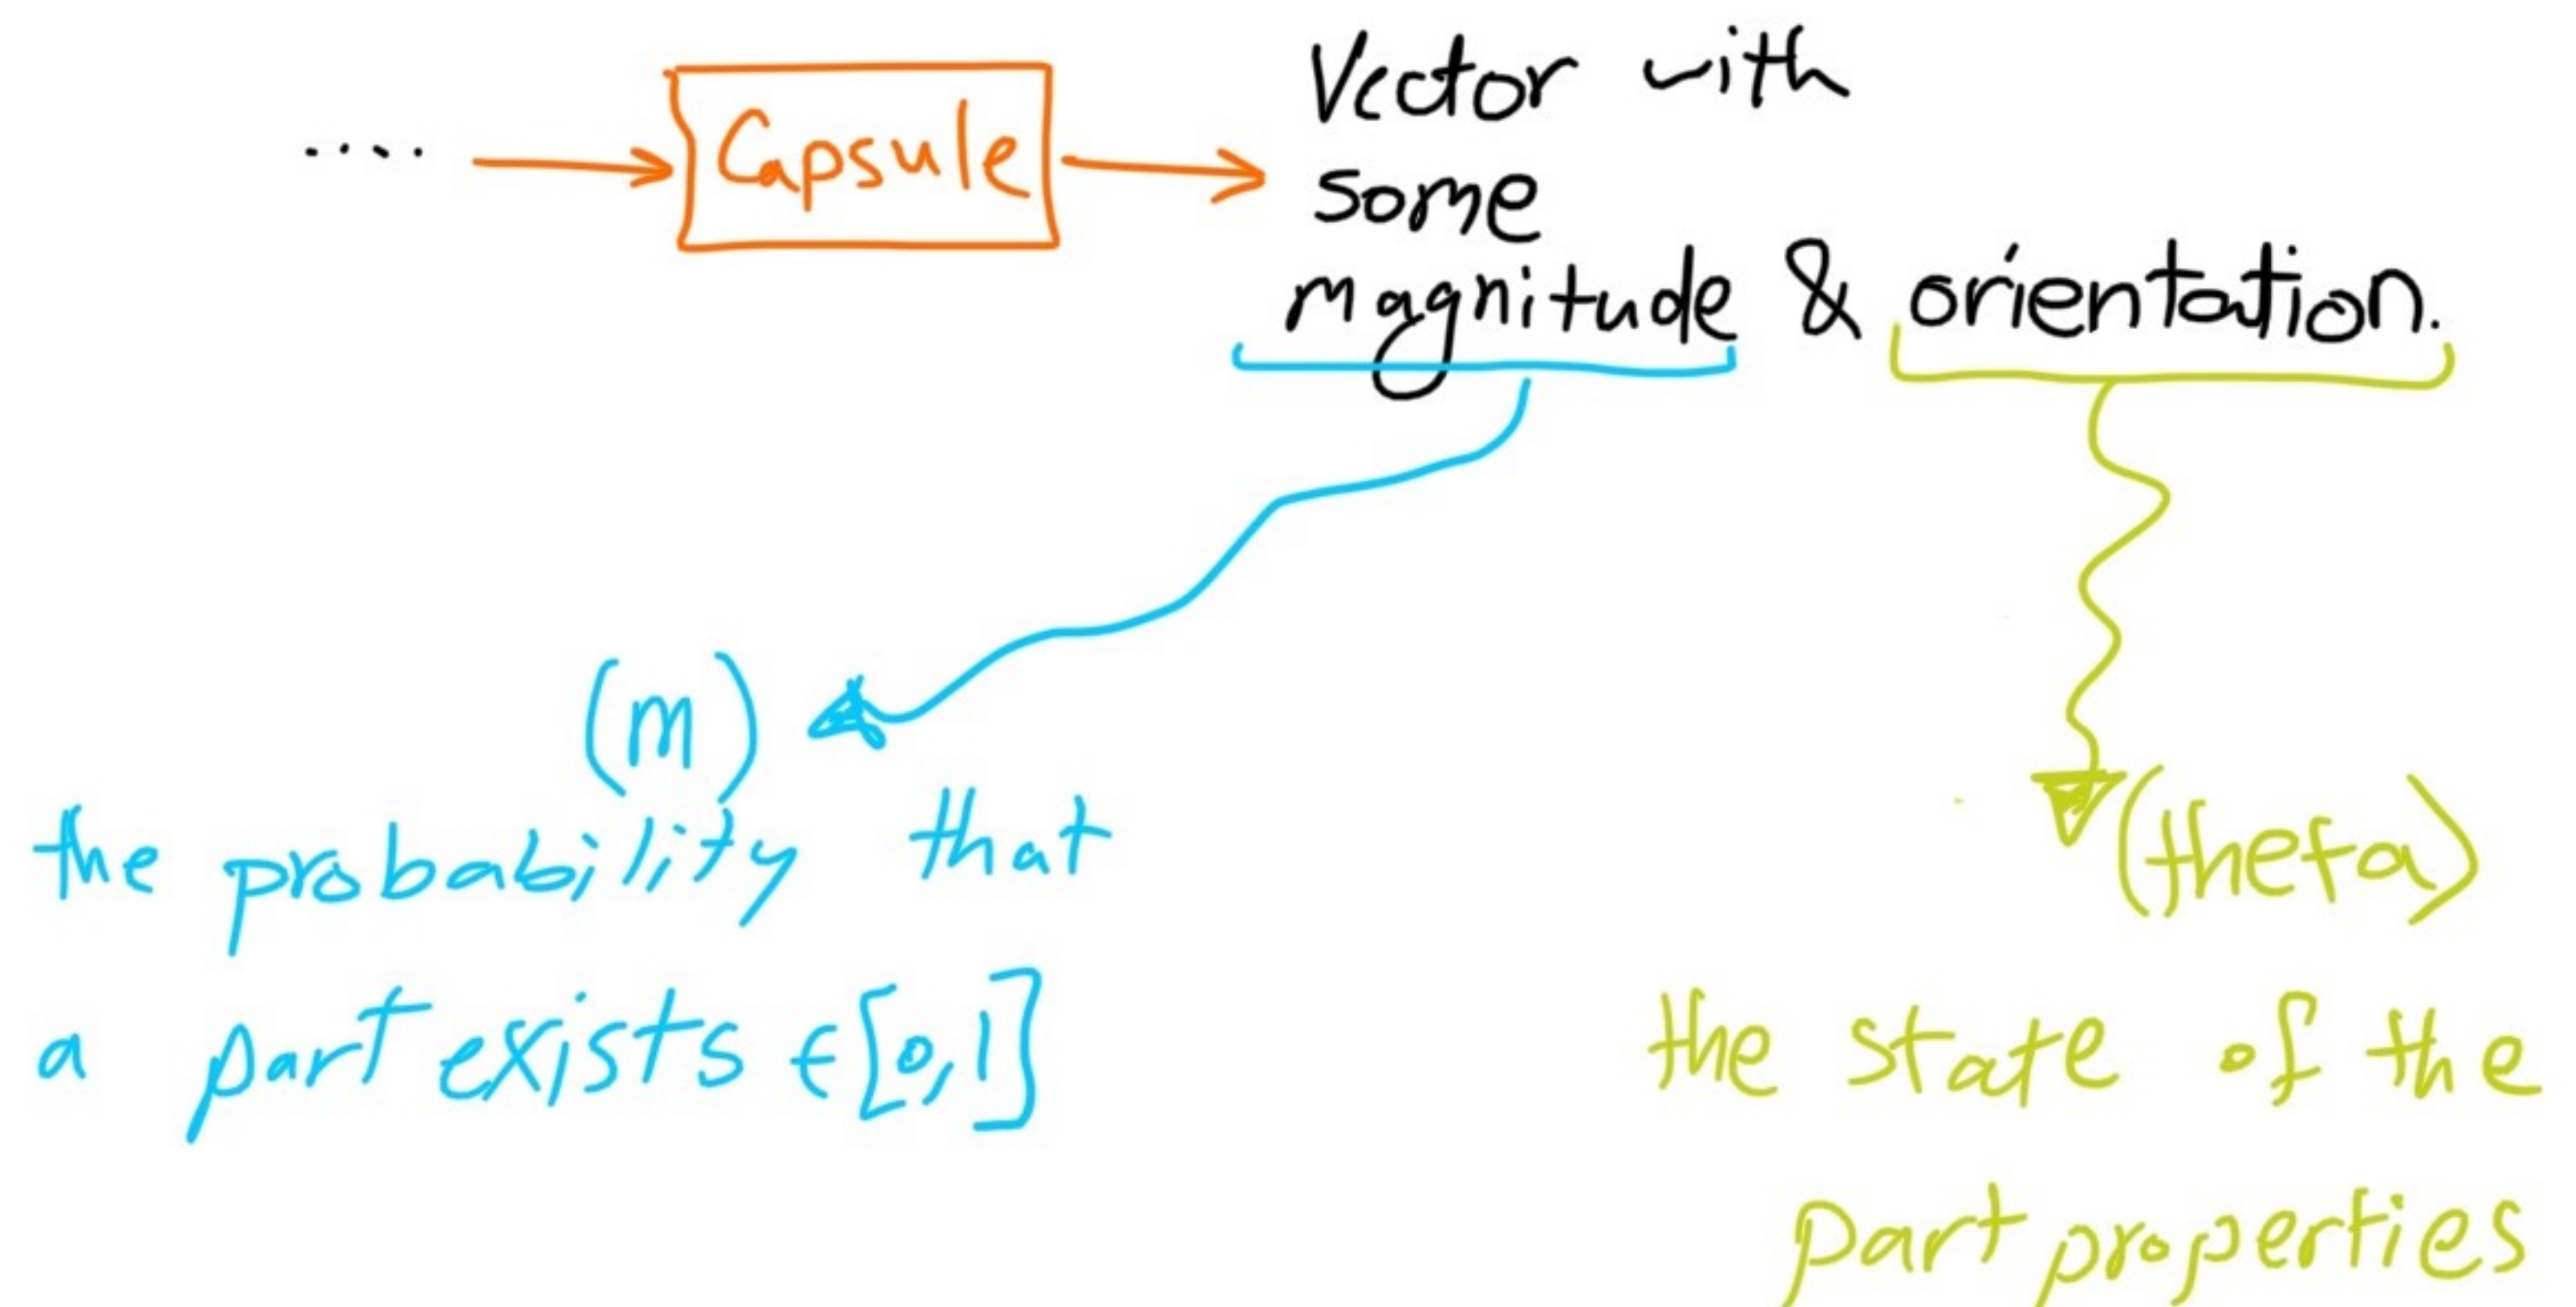
pooling throws away some info,
i.e., discards some pixels.

In classification, pooling is fine.
But in something like face
recognition it is bad.

Instead, some people use strategies
to keep some spatial info. An example
is **Capsule Network**, which learns spatial
relationships between parts.



Capsule: Collection of nodes that contain info about a sub-part of the image (width, orientation, color, ...)



The output of these capsules can be used to make a "tree".

