

CIFAR problem Using PyTorch

to check if CUDA is available:

```
train_on_gpu = torch.cuda.is_available()
```

Load the data

```
transform = transforms.Compose([  
    transforms.ToTensor(),  
    transforms.Normalize((1/2, 1/2, 1/2), (1/2, 1/2, 1/2))  
])
```

```
train_data = datasets.CIFAR10('data', train=True,  
                               download=True,  
                               transform=transform)
```

↑
torchvision
↓

```
test_data = datasets.CIFAR10('data', train=False,  
                               download=True,  
                               transform=transform)
```

data to validate {

```
num_train = len(train_data)  
indices = list(range(num_train))  
np.random.shuffle(indices)  
split = int(np.floor(valid_size * num_train))  
train_idx, valid_idx = indices[split:], indices[:split]
```

→ 0.2

to make batches

```
train_sampler = SubsetRandomSampler(train_idx)
valid_sampler = SubsetRandomSampler(valid_idx)

train_loader = torch.utils.data.DataLoader(train_data,
                                             batch_size=batch_size,
                                             sampler=train_sampler,
                                             num_workers=nw)

valid_loader = torch.utils.data.DataLoader(valid_data,
                                             batch_size=batch_size,
                                             sampler=valid_sampler,
                                             num_workers=nw)

test_loader = torch.utils.data.DataLoader(test_data,
                                           batch_size=batch_size,
                                           num_workers=nw)
```

it is a classification task

```
classes = ['airplane', ..., 'truck']
```


model

```
class Net(nn.Module):
```

```
    def __init__(self):
```

```
        super(Net, self).__init__()
```

```
    sees 32x32x3 ← self.conv1 = nn.Conv2d(3, 16, 3, padding=1)
```

```
    sees 16x16x16 ← self.conv2 = nn.Conv2d(16, 32, 3, padding=1)
```

```
    sees 8x8x32 ← self.conv3 = nn.Conv2d(32, 64, 3, padding=1)
```

```
        self.pool = nn.MaxPool2d(2, 2)
```

```
        self.fc1 = nn.Linear(64 * 4 * 4, 500)
```

```
        self.fc2 = nn.Linear(500, 10)
```

```
        self.dropout = nn.Dropout(.25)
```



```
def forward(self, x):
```

```
    x = self.pool(F.relu(self.conv1(x)))
```

```
    x = self.pool(F.relu(self.conv2(x)))
```

```
    x = self.pool(F.relu(self.conv3(x)))
```

```
    x = x.view(-1, 64 * 4 * 4)
```

```
    x = self.dropout(x)
```

```
    x = F.relu(self.fc1(x))
```

```
    x = self.dropout(x)
```

```
    x = self.fc2(x)
```

```
    return x
```

Use the defined model

```
model = Net()
```

```
if train_on_gpu: model.cuda()
```

move tensors to Cuda

Training step

```
Criterion = nn.CrossEntropyLoss()
```

```
optimizer = optim.Adam(model.parameters(), lr=0.001)
```

```
valid_loss_min = np.Inf
```


for epoch in range(epochs):

train_loss = valid_loss = 0.0

model.train() \leftarrow train mode

for data, target in train_loader:

if train_on_gpu:

data, target = data.cuda(), target.cuda()

optimizer.zero_grad()

output = model(data)

loss = criterion(output, target)

loss.backward()

optimizer.step()

train_loss += loss.item() * data.size(0)

model.eval() \leftarrow eval mode

for data, target in valid_loader:

if train_on_gpu:

data, target = data.cuda(), target.cuda()

output = model(data)

loss = criterion(output, target)

valid_loss += loss.item() * data.size(0)

Compute the final values for losses

train_loss /= len(train_loader.dataset)

valid_loss /= len(valid_loader.dataset)

Save only if there is an improvement

if valid_loss <= valid_loss_min:

torch.save(model.state_dict(), 'model.pt')

valid_loss_min = valid_loss

Summary of the Common Strategy:

1. After each conv layer, we do relu and then maxpooling.
2. Then, we flatten the data.
3. Pass through dropouts