


Auto encoders



In PyTorch: option 1: MLP

* the output image and input image are

data handling as before

kinda
different.
(blurry)

```
class AutoEncoder(nn.Module):
```

```
    def __init__(self, encoding_dim):
```

```
        super(AutoEncoder, self).__init__()
```

```
    # encoder:
```

```
        self.fc1 = nn.Linear(28*28, encoding_dim)
```

```
    # decoder
```


```
        self.fc2 = nn.Linear(encoding_dim, 28*28)
```

```
    def forward(self, x):
```

```
        x = F.relu(self.fc1(x))
```

```
        x = F.sigmoid(self.fc2(x))
```

```
        return x
```



encoding - dim = 32


model = AutoEncoder(encoding - dim)

★ criterion = nn.MSELoss() → good for pixels comparison
not probabilities
optimizer = optim.Adam(model.parameters(), lr = 0.001)

train loop as before; but we are not interested in labels here. We just want to compare the input image with reconstructed image.


↳ at each step: loss = criterion(outputs, images)
this tells us how good of a reconstruction our model does

option 2: conv. layers ✖ a better solution but still has issue of "blocks" in the image.
key point: transpose conv. layers




→ we take each pixel, and use a kernel for upsampling. e.g., we use a 3x3 kernel for a 2x2 image. (let's call this A)

For each pixel in the image, we multiply the value of it with each element of the kernel.



Convolutional Autoencoders



In Pytorch, transpose Conv Layers are `nn.ConvTranspose2d`.

```
class ConvAutoencoder(nn.Module):
```

```
    def __init__(self):
```

```
        super(ConvAutoencoder, self).__init__()
```

```
        # encoder layers
```

```
depth: 1 → 16 self.conv1 = nn.Conv2d(1, 16, 3, padding=1)
```

```
depth: 16 → 4 self.conv2 = nn.Conv2d(16, 4, 3, padding=1)
```

```
self.pool = nn.MaxPool2d(2, 2)
```

```
        # decoder layers
```

```
self.t_conv1 = nn.ConvTranspose2d(4, 16, 2, stride=2)
```

```
self.t_conv2 = nn.ConvTranspose2d(16, 1, 2, stride=2)
```

```
    def forward(self, x):
```

```
        x = F.relu(self.conv1(x))
```

```
        x = self.pool(x)
```

```
        x = F.relu(self.conv2(x))
```

```
        x = self.pool(x)
```

```
        x = F.relu(self.t_conv1(x))
```

```
        x = F.Sigmoid(self.t_conv2(x))
```

```
        return x
```


model = ConvAutoencoder()

criterion = nn.MSELoss()

most of the train loop is as before.
just note following statements:

outputs = model(images)

loss = Criterion(outputs, images)

i.e., we want outputs and images be ideally identical.

option 3: upsampling

smoother and
more similar images

We change option 2's decoder by adding a few layers. (upsampling)

*idea: replace a transpose Conv. layer with a conv. layer + an upsample layer.

The change with option 2's code is just in the "forward" portion of the module.

encoder
:

decoder


$x = F.\text{upsample}(x, \text{scale_factor}=2, \text{mode}='nearest')$

$x = F.\text{relu}(\text{self.conv4}(x))$

$x = F.\text{upsample}(x, \text{scale_factor}=2, \text{mode}='nearest')$

$x = F.\text{sigmoid}(\text{self.conv5}(x))$

Denoising Autoencoders



In PyTorch

```
class ConvDenoiser(nn.Module):
```

```
    def __init__(self):
```

```
        super(ConvDenoiser, self).__init__()
```

Encoder

```
self.Conv1 = nn.Conv2d(1, 32, 3, padding=1)
```

```
self.Conv2 = nn.Conv2d(32, 16, 3, padding=1)
```

```
self.Conv3 = nn.Conv2d(16, 8, 3, padding=1)
```

```
self.Pool = nn.MaxPool2d(2, 2)
```

Decoder

```
self.t_Conv1 = nn.ConvTranspose2d(8, 8, 3, stride=2)
```

```
self.t_Conv2 = nn.ConvTranspose2d(8, 16, 2, stride=2)
```

```
self.t_Conv3 = nn.ConvTranspose2d(16, 32, 2, stride=2)
```

```
self.Conv_out = nn.Conv2d(32, 1, 3, padding=1)
```



```
def forward (self, x):
```

```
    x = F.relu(self.conv1(x))
```

```
    x = self.pool(x)
```

```
    x = self.relu(self.conv2(x))
```

```
    x = self.pool(x)
```

```
    x = F.relu(self.conv3(x))
```

```
    x = self.pool(x)
```

```
    x = F.relu(self.t_conv1(x))
```

```
    x = F.relu(self.t_conv2(x))
```

```
    x = F.relu(self.t_conv3(x))
```

```
    x = F.sigmoid(self.conv_out(x))
```

```
    return x
```