






# Implementing style transfer in images



$$\underbrace{L_{\text{Content}}}_{\text{Content loss}} = \frac{1}{2} \sum (T_c - C_c)^2$$



# Gram Matrix



List of Gram Matrices of style image ( $S_s$ )  
List of Gram Matrices of target image ( $T_s$ )

$$\underbrace{L_{\text{style}}}_{\text{style loss}} = a \sum_i w_i (T_{s,i} - S_{s,i})^2$$

$$\text{Total loss} = \underbrace{\alpha}_{\text{content weight}} L_{\text{content}} + \underbrace{\beta}_{\text{style weight}} L_{\text{style}}$$



- Style weight is usually much larger than the content weight.
- to add more style to target image, decrease  $\frac{\alpha}{\beta}$ .

In PyTorch:

from PIL import Image → load any image

from io import BytesIO

import numpy as np

import torch

import torch.optim as optim

import requests

from torchvision import transforms, models

# get feature layers of vgg19: Conv. and pooling layers  
↑

vgg = models.vgg19(pretrained=True).features

# freeze the parameters:

for param in vgg.parameters():

param.requires\_grad\_(False)



# move to device

```
device = torch.device("cuda" if torch.cuda.is_available()
                      else "cpu")
```

```
vgg.to(device)
```

# Load Content/Style images

```
def load_image (img_path, max_size = 400,
                shape = None):
```

```
    if "http" in img_path:
```

```
        response = requests.get(img_path)
```

```
        image = Image.open(BytesIO(response.content))
```

```
        • Convert('RGB')
```

```
    else:
```

```
        image = Image.open(img_path).convert('RGB')
```

actual  
loading



resize  
all to  
a  
managable  
size

```
if max(image.size) > max_size:
```

```
    size = max_size
```

```
else:
```

```
    size = max(image.size)
```

```
if shape is not None:
```

```
    size = shape
```

```
in_transform = transforms.Compose([  
    transforms.Resize(size),  
    transforms.ToTensor(),  
    transforms.Normalize(  
        (0.485, 0.456, 0.406),  
        (0.229, 0.224, 0.225))])
```

discard  
alpha  
channel

```
→ image = in_transform(image)[:3, :, :].unsqueeze(0)  
return image
```

```
content = load_image('img1.jpg').to(device)
```

```
style = load_image('img2.jpg', shape = content.shape[-2:]).  
                                             to(device)
```

to align easier



un-  
norm-  
alize  
an  
image

```
def im_Convert(tensor):
```

```
    image = tensor.to('cpu').clone().detach()
```

```
    image = image.numpy().squeeze()
```

```
    image = image.transpose(1,2,0)
```

```
    image = image * np.array((0.229, 0.224, 0.225))  
        + np.array((0.485, 0.456, 0.406))
```

```
    image = image.clip(0,1)
```

```
    return image
```

```
fig, (ax1, ax2) = plt.subplots(1,2, figsize=(20,10))
```

```
ax1.imshow(im_Convert(content))
```

```
ax2.imshow(im_Convert(style))
```

→ 

```
def get_features(image, model, layers=None):
```

```
    if layers is None:
```

```
        layers = { '0' : 'Conv1-1',
```

```
                  '5' : 'Conv2-1',
```

```
                  '10' : 'Conv3-1',
```

```
                  '19' : 'Conv4-1',
```

```
                  '21' : 'Conv4-2',
```

```
                  '28' : 'Conv5-1' }
```

pass  
image  
through  
layers



features = {}

x = image

Contains all the modules of a model

for name, layer in model.\_modules.items():

    x = layer(x)

    if name in layers:

        features[layers[name]] = x

return features

def gram\_matrix(tensor):

    b, d, h, w = tensor.size()

    tensor = tensor.view(b\*d, h\*w)

    gram = torch.mm(tensor, tensor.t())

    return gram

# finally!

content\_features = get\_features(content, vgg)

style\_features = get\_features(style, vgg)

style\_grams = {layer: gram\_matrix(style\_features[layer])

for layer in style\_features}

better initialize with content img  
↓

target = content.clone().requires\_grad\_(True).to(device)



style\_weights = { 'Conv1-1' : 1.0,  
                  'Conv2-1' : 0.75,  
                  'Conv3-1' : 0.2,  
                  'Conv4-1' : 0.2,  
                  'Conv5-1' : 0.2 }

content\_weight = 1  $\rightarrow$  alpha

style\_weight = 1e6  $\rightarrow$  beta

optimizer = optim.Adam([target], lr=0.003)

for \_ in range(steps):

    target\_features = get\_features(target, vgg)

    content\_loss = torch.mean(  
                                (target\_features['conv4-2']  
                                - content\_features['conv4-2'])  
                                \*\* 2)



Style\_loss = 0

for layer in style\_weights:

- target\_features = target\_features[layer]
- target\_gram = gram\_matrix(target\_feature)
- —, d, h, w = target\_feature.shape
- style\_gram = style\_grams[layer]
- layer\_style\_loss =  
    style\_weights[layer] \*  
    torch.mean((target\_gram -  
                  style\_gram)\*\*2)
- style\_loss += layer\_style\_loss / (d \* h \* w)

- total\_loss = Content\_weight \*  
    Content\_loss +  
    style\_weight \*  
    style\_loss

update  
the  
target  
image

- optimizer.zero\_grad()
- total\_loss.backward()
- optimizer.step()

\* target is our final output.