

# Implement RNN in PyTorch

★ (a simple time-series prediction) ★

## # Initial Reqs

```
from torch import nn
import numpy as np
import matplotlib.pyplot as plt
```

## # stuff for visualizing the data

```
plt.figure(figsize=(8, 5))
seq_length = 20 → # of data points in each batch
time_steps = np.linspace(0, np.pi, seq_length + 1)
data = np.sin(time_steps)
data.resize((seq_length + 1, 1)) → adds a dimension
```

```
x = data[:, -1]
```

```
y = data[1:]
```

```
plt.plot(time_steps[1:], x, 'r.', label='input, x')
```

```
plt.plot(time_steps[1:], y, 'b.', label='target, y')
```

```
plt.legend(loc='best')
```

```
plt.show()
```



## # Define RNN

```
class RNN (nn.Module):  
    def __init__(self, input_size, output_size,  
                  hidden_dim, n_layers):  
        super(RNN, self).__init__()  
        self.hidden_dim = hidden_dim
```

Actual RNN!  $\leftarrow$  self.rnn = nn.RNN(input\_size, hidden\_dim,  
 n\_layers, batch\_first=True)  
self.fc = nn.Linear(hidden\_dim, output\_size)

batch size  $\times$  seq-len  $\times$  input\_size  $\rightarrow$  def forward(self,  $x$ , hidden):  $\rightarrow$  hidden state  
n\_layers  $\times$  batch\_size  $\times$  hidden\_dim

$r\_out, hidden = self.rnn(x, hidden)$

RNN output

batch size  $\times$  time-step  $\times$  hidden\_dim

$r\_out = r\_out.view(-1, self.hidden\_dim)$

new shape:

batch\_size \* timeStep  $\times$  hidden dim

We are doing some sort of flattening for FC layer.

output = self.fc(r\_out)

return output, hidden



# Let's test the RNN!

```
my_rnn = RNN(input_size=1, output_size=1,  
             hidden_dim=10, n_layers=2)
```

```
time_steps = np.linspace(0, np.pi, seq_length)
```

```
data = np.sin(time_steps)
```

```
data.resize((seq_length, 1))
```

add a dim  
for batch.

```
test_input = torch.Tensor(data).unsqueeze(0)
```

```
out, h = my_rnn(test_input, None)
```

# Training our RNN

```
input_size = 1
```

```
output_size = 1
```

```
hidden_dim = 32
```

```
n_layers = 1
```

} hyperparams

```
rnn = RNN(input_size, output_size,  
         hidden_dim, n_layers)
```

Loss  
&  
Opt.

```
Criterion = nn.MSELoss()
```

```
optimizer = torch.optim.Adam(rnn.parameters(), lr=.01)
```

standard for RNNs



## # Training Loop

```
def train(rnn, n-steps, print-every):  
    hidden = None
```

```
    for i, step in enumerate(range(n-steps)):
```

```
        time-steps = np.linspace(step * np.pi,  
                                (step+1) * np.pi,  
                                seq-length + 1)
```

```
        data = np.sin(time-steps)
```

```
        data.resize((seq-length + 1, 1))
```

```
        x = data[:-1]
```

```
        y = data[1:]
```

```
        x-tensor = torch.Tensor(x).unsqueeze(0)
```

```
        y-tensor = torch.Tensor(y)
```

```
        prediction, hidden = rnn(x-tensor, hidden)
```

```
# ~ ~ ~ ~ ~ ~ ~ ~ #
```

```
# ~ representing memory ~ #
```

```
    ↖ hidden = hidden.data
```

avoid backpropagating through the entire history

```
    loss = criterion(prediction, y-tensor)
```



optimizer.zero\_grad()

loss.backward()

optimizer.step()

# Some fancy visualization

if i % print\_every == 0:

plt.plot(time\_steps[1:], x, 'r')

plt.plot(time\_steps[1:],

prediction.data.numpy().flatten(),  
'b')

plt.show()

n\_steps = 75

print\_every = 15

trained\_rnn = train(rnn, n\_steps, print\_every)