



An Exercise in LSTM — (PyTorch)

* learning about a text one character at a time and generating new text one character at a time (character-wise RNN)

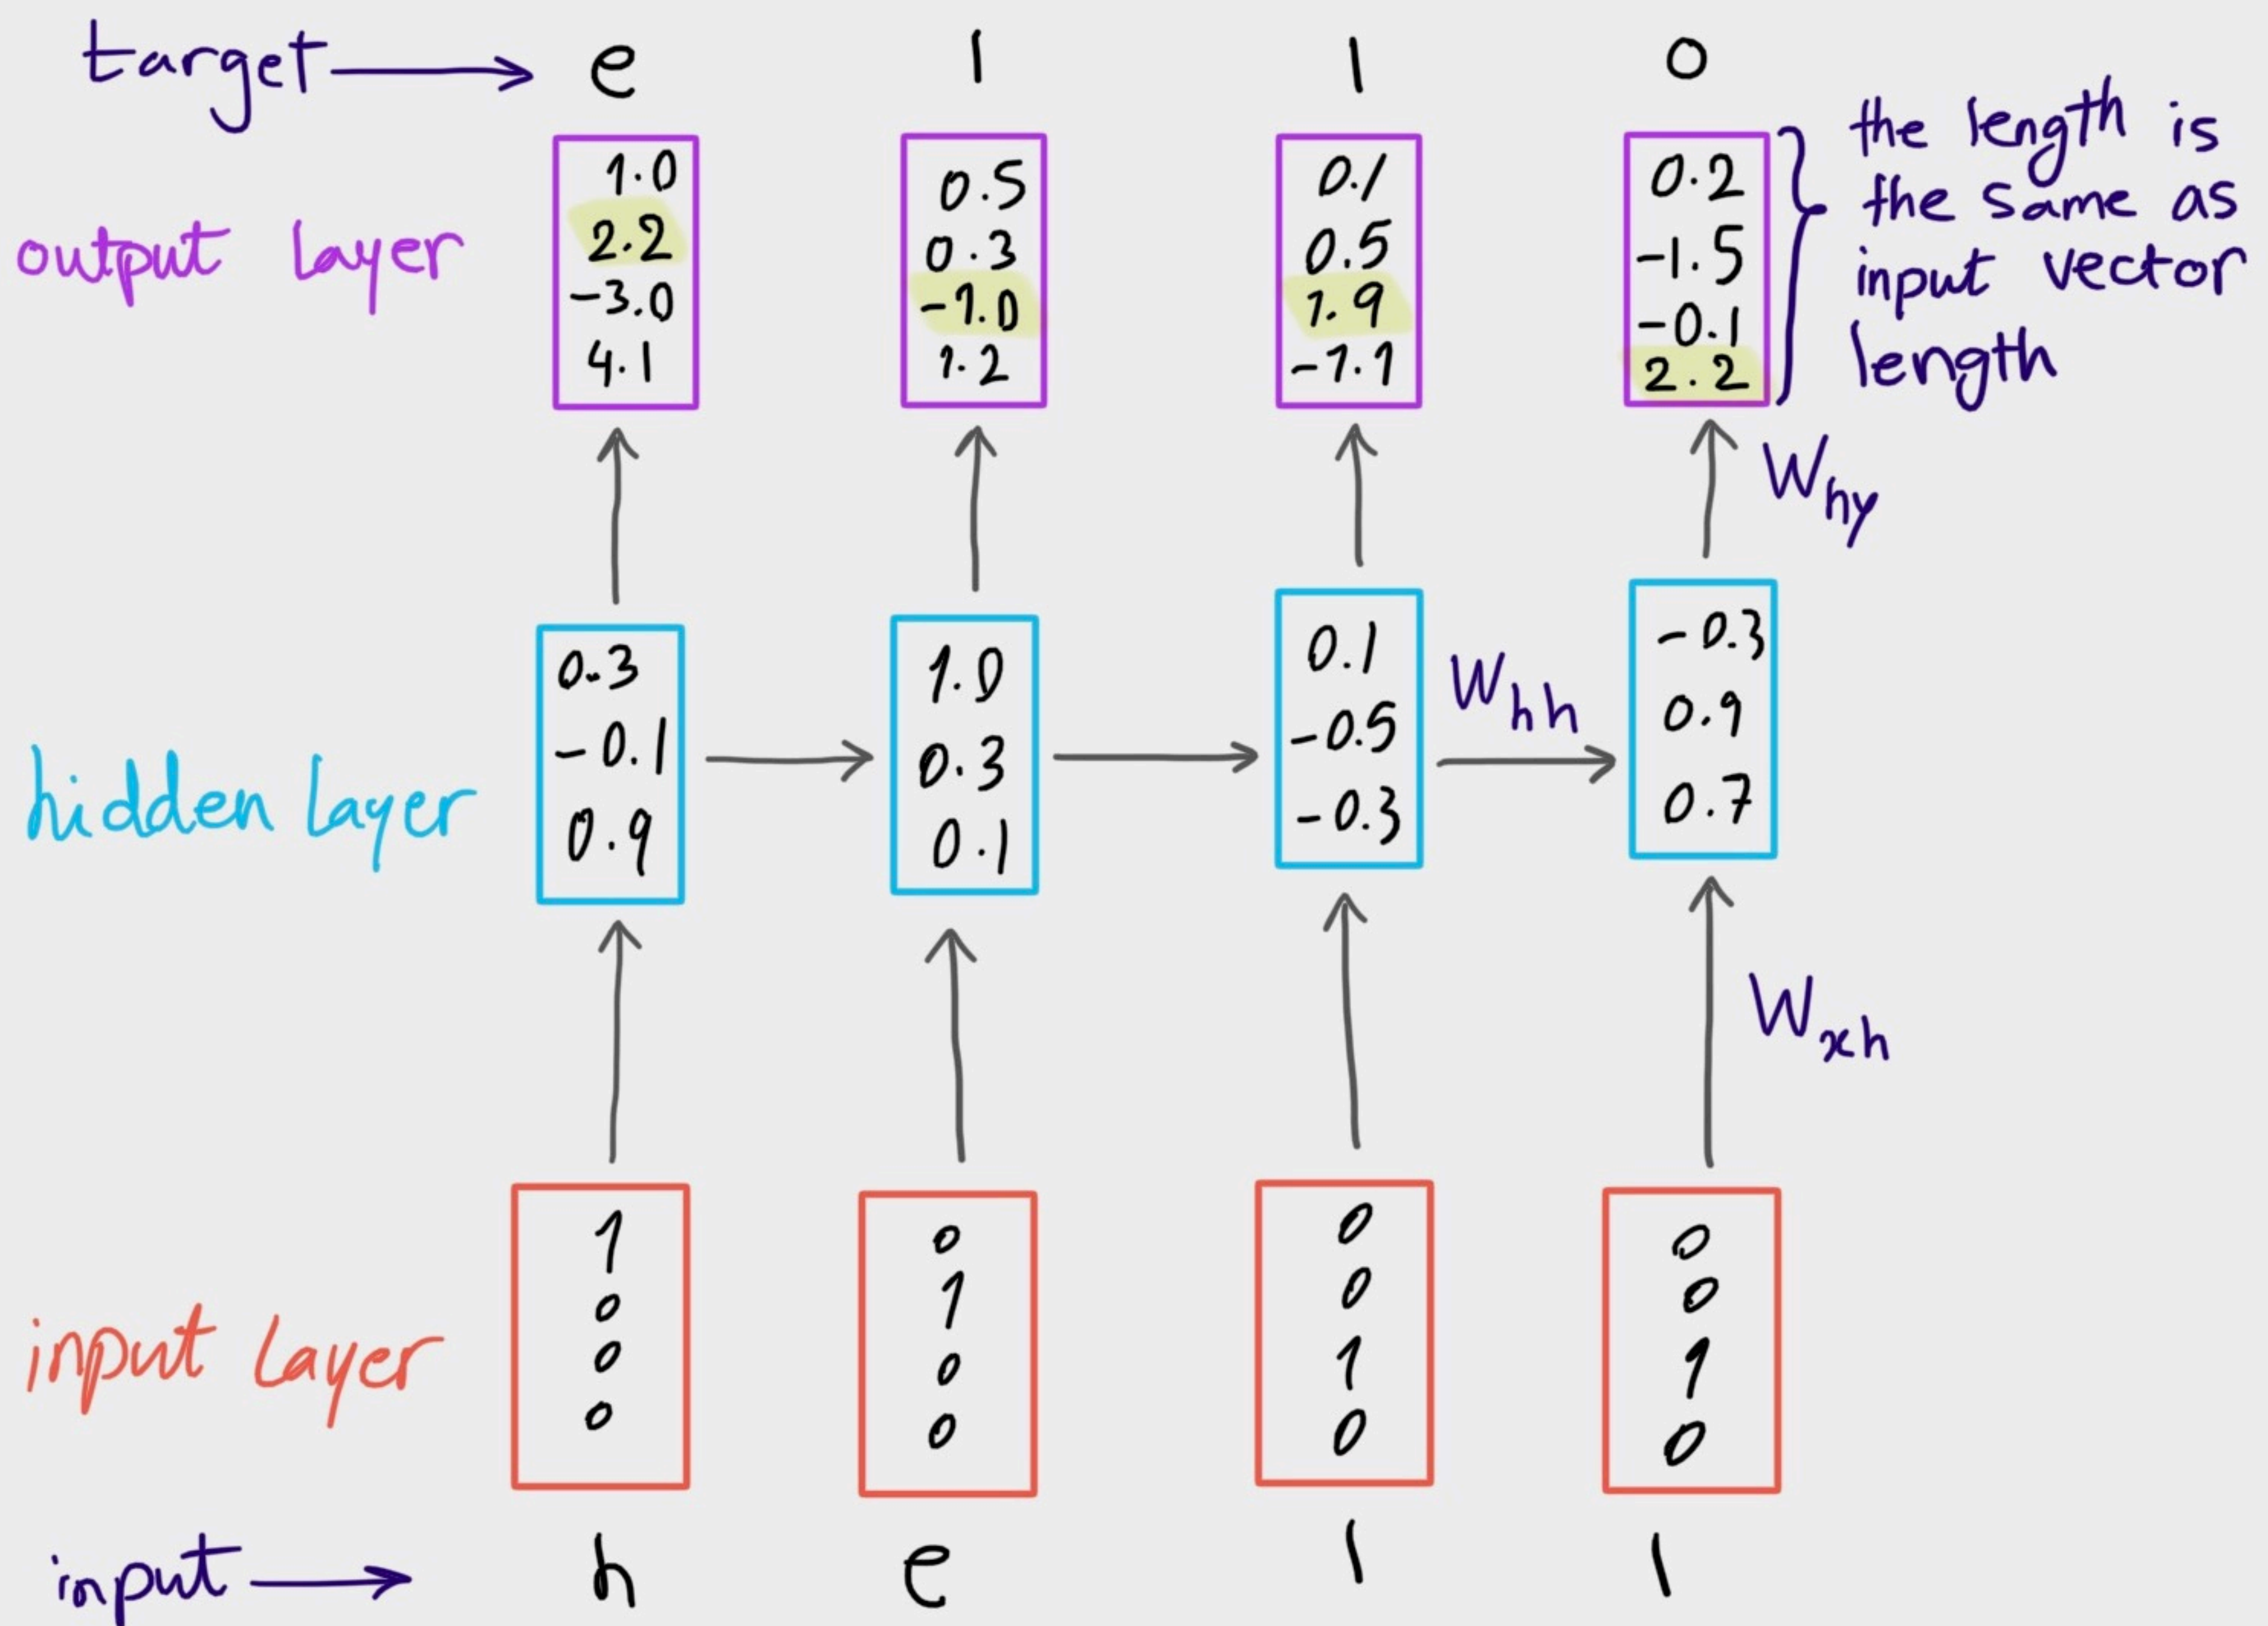


- We should 'batch' the sequences to better use matrix mult capabilities.

[● ▲ ◆ ■ ✕ + ♥ ♠ 🍎 🍷 ÷ %]



Now in PyTorch what we want to implement:



Code

```
import numpy as np
import torch
from torch import nn
import torch.nn.functional as F
```


Loading the data

with open('file.txt', 'r') as f:

```
text = f.read()
```

Tokenization

```
chars = tuple(set(text))
```

```
int2char = dict(enumerate(chars))
```

```
char2int = {ch: ii for ii, ch in int2char.items()}
```

```
encoded = np.array([char2int[ch] for ch in text])
```

Pre-processing

```
def one-hot(arr, n_labels):
```

```
    oh = np.zeros((np.multiply(*arr.shape), n_labels),  
                  dtype=np.float32)
```

```
    oh[np.arange(oh.shape[0]), arr.flatten()] = 1.0
```

```
    oh = oh.reshape((*arr.shape, n_labels))
```

```
    return oh
```


Making Batches

```
def get_batches (arr, batch_size, seq_length):
```

```
    n_batches = len(arr) // (batch_size * seq_len)
```

```
    arr = arr[: n_batches * batch_size * seq_len]
```

```
    arr = arr.reshape((batch_size, -1))
```

```
    for n in range(0, arr.shape[1], seq_length):
```

```
        x = arr[:, n:n+seq_length]
```

all rows → *Some columns*

```
        y = np.zeros_like(x)
```

```
        try:
```

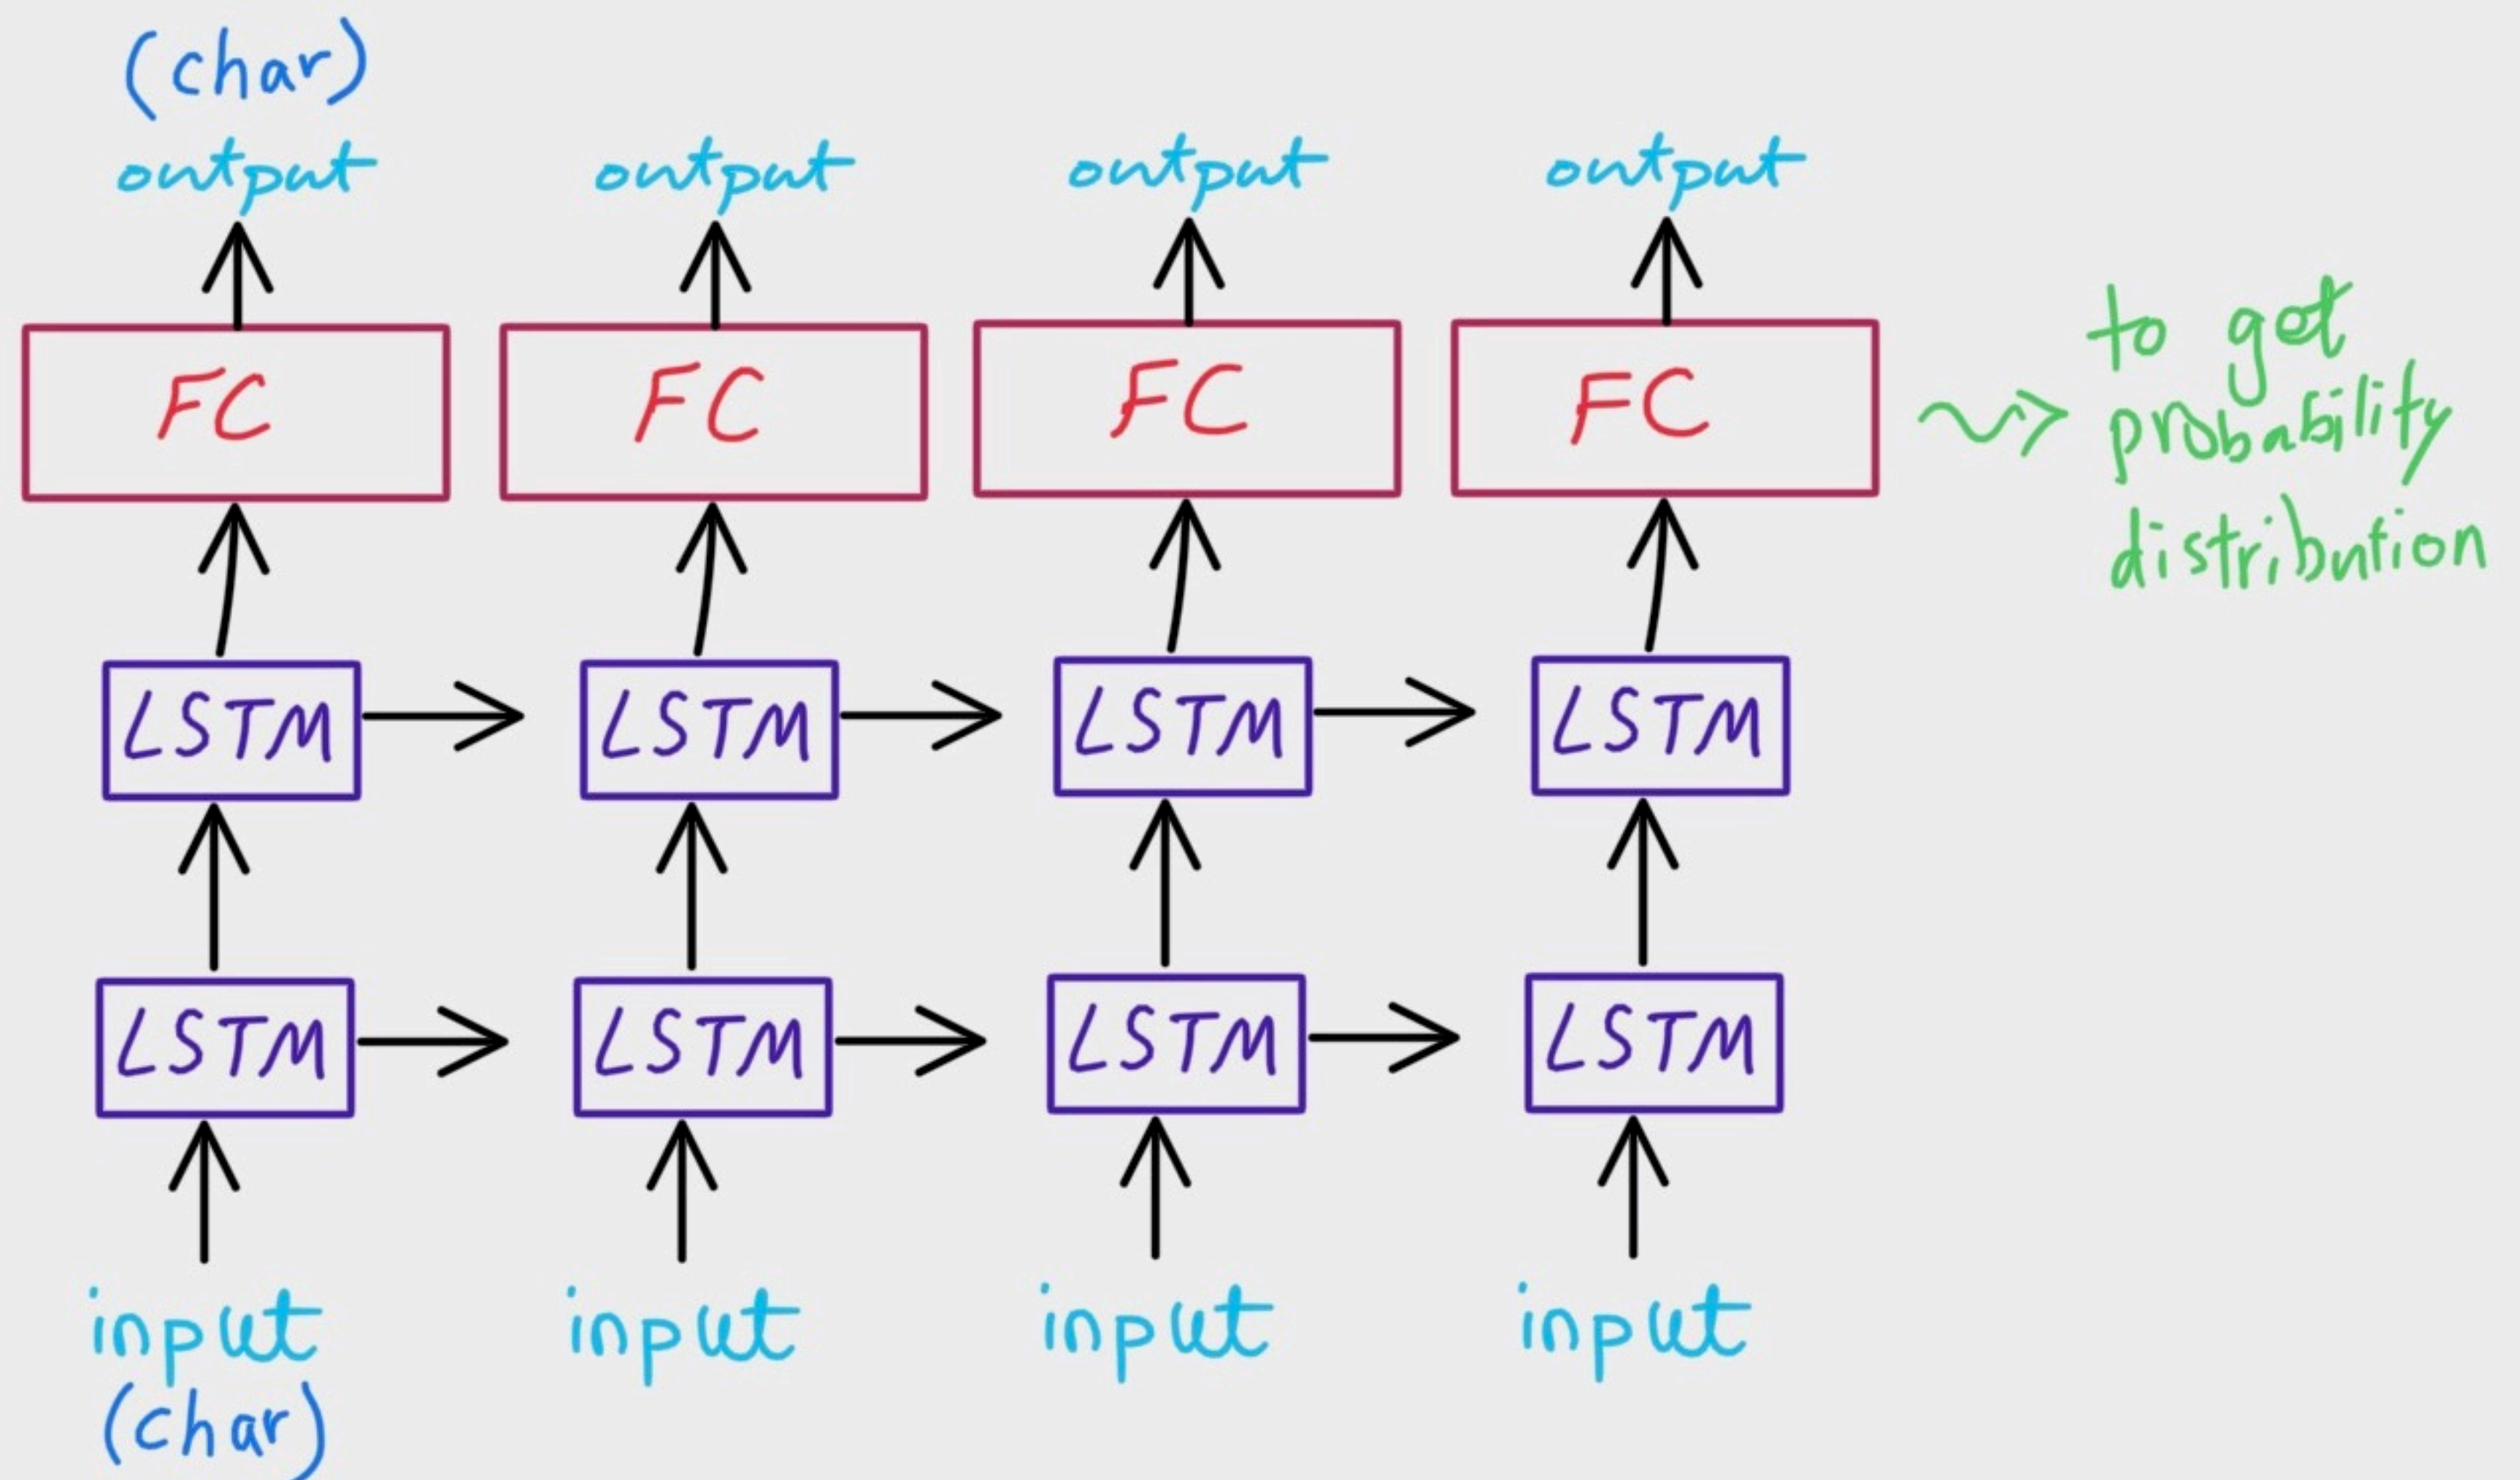
```
            y[:, :-1], y[:, -1] = x[:, 1:], arr[:, n+seq_length]
```

```
        except IndexError:
```

```
            y[:, :-1], y[:, -1] = x[:, 1:], arr[:, 0]
```

```
        yield x, y
```


The model



```
class CharRNN(nn.Module):
```

```
    def __init__(self, tokens, n_hidden=256, n_layers=2,  
                  drop_prob=0.5, lr=0.001):
```

```
        super().__init__()
```

```
        self.drop_prob = drop_prob
```

```
        self.n_layers = n_layers
```

```
        self.n_hidden = n_hidden
```

```
        self.lr = lr
```

Making the Dictionary

```
{ self.chars = tokens  
  self.int2char = dict(enumerate(self.chars))  
  self.char2int = {ch: ii for ii, ch in self.int2char.items()}
```


LSTM Layer in PyTorch

```
nn.LSTM(input_size, n_hidden, n_layers,  
        dropout=drop_prob, batch_first=True)
```

layers {

```
self.lstm = nn.LSTM(len(self.chars), n_hidden,  
                    n_layers, dropout=drop_prob,  
                    batch_first=True)
```

→ dropout between LSTM layers

```
self.dropout = nn.Dropout(drop_prob)  
self.fc = nn.Linear(n_hidden, len(self.chars))
```

```
def forward(self, x, hidden):  
    r_output, hidden = self.lstm(x, hidden)  
    out = self.dropout(r_output)  
    out = out.view(-1, self.n_hidden) → stacking LSTM hidden outputs  
  
    out = self.fc(out)  
    return out, hidden
```



```

def init_hidden(self, batch-size):
    w = next(self.parameters()).data
    if train_on_gpu:
        hidden = (w.new(self.n_layers, batch-size, self.n-hidden).
                    zero_().cuda(),
                  w.new(self.n_layers, batch-size, self.n-hidden).
                    zero_().cuda())
    else:
        hidden = (w.new(self.n_layers, batch-size, self.n-hidden).
                    zero_(),
                  w.new(self.n_layers, batch-size, self.n-hidden).
                    zero_())
    return hidden

```

Training

```
def train(model, data, n_epochs=2, batch_size=32,  
          seq_length=128, lr=0.001, clip=5,  
          val_frac=0.1):
```

fraction of
holdout data for validation

gradient
clipping

```
    model.train()
```

```
    opt = torch.optim.Adam(model.parameters(), lr=lr)
```

```
    criterion = nn.CrossEntropyLoss()
```

```
    val_idx = int(len(data) * (1 - val_frac))
```

```
    data, val_data = data[:val_idx], data[val_idx:]
```

```
    if train_on_gpu:
```

```
        model.cuda()
```

```
    counter = 0
```

```
    n_chars = len(model.chars)
```



```
for e in range(epochs):
```

```
    h = model.init_hidden(batch_size)
```

```
    for x, y in get_batches(data, batch_size, seq_length):
```

```
        counter += 1
```

```
        x = one_hot(x, n_chars)
```


```
        input, targets = torch.from_numpy(x),  
                           torch.from_numpy(y)
```

```
        if train_on_gpu:
```

```
            inputs, targets = inputs.cuda(), targets.cuda()
```

```
        h = tuple([each.data for each in h])
```

create new vars
for hidden state
to avoid backprop
on the entire history



```
        model.zero_grad()
```

```
        output, h = model(inputs, h)
```

```
        loss = Criterion(output, targets.view(batch_size *  
                                              seq_length))
```

```
        loss.backward()
```

prevent
gradient
exploding → nn.utils.clip_grad_norm_(model.parameters, clip)

opt.step()

if counter % print_every == 0:

val_h = model.init_hidden(batch_size)

losses = []

model.eval()

for x, y in get_batches(val_data, batch_size,
seq_length):

x = one_hot(x, n_chars)

x, y = torch.from_numpy(x),
torch.from_numpy(y)

val_h = tuple([each.data for each in val_h])

inputs, targets = x, y

if train_on_gpu:

inputs, targets = inputs.cuda(), targets.cuda()

out, val_h = model(inputs, val_h)

loss = criterion(out, targets.view(batch_size *
seq_length))

losses.append(loss.item())

model.train()

Let's train now!

n-hidden = 512

n-layers = 2

model = CharRNN(chars, n-hidden, n-layers)

batch-size = 128

seq-length = 100

epochs = 2

train(model, encoded, epochs=epochs, batch-size=batch-size,
seq-length=seq-length, lr=0.001)

Save the Model

model_name = f'model-{{epoch}}-{{n-hidden}}-{{n-layers}}'

checkpoint = {'n-hidden': model.n-hidden,
'n-layers': model.n-layers,
'state-dict': model.state_dict(),
'tokens': model.chars}

with open(model_name, 'wb') as f:

torch.save(checkpoint, f)

Use the trained Model

```
def predict(model, char, h=None, top_k=None):
```

make
tensor

```
    x = np.array([net.char2int[char]])  
    x = one_hot(x, len(net.chars))  
    inputs = torch.from_numpy(x)
```

```
    if train_on_gpu:
```

```
        inputs = inputs.cuda()
```

detach
hidden
state
from
history

```
    → h = tuple([each.data for each in h])  
    out, h = Model(inputs, h)
```

```
    p = F.softmax(out, dim=1).data
```

```
    if train_on_gpu:
```

```
        p = p.cpu()
```

```
    if top_k is None:
```

```
        top_ch = np.arange(len(model.chars))
```

```
    else:
```

```
        p, top_ch = p.topk(top_k)
```

```
        top_ch = top_ch.numpy().squeeze()
```

only consider
if the predicted int
is in dict. ↙

select
the next
likely
character

```
    p = p.numpy().squeeze()
```

```
    char = np.random.choice(top_ch, p=p/p.sum())
```



```
return net.int2char(char), h
```

How to use prediction to do text generation?

```
def sample(model, size, prime='The', top_k=None):  
    if train_on_gpu:
```

```
        model.cuda()
```

```
    else:
```

```
        model.cpu()
```

```
    chars = [ch for ch in prime]
```

```
    h = model.init_hidden(1)
```

```
    for ch in prime:
```

```
        char, h = predict(model, ch, h, top_k=top_k)
```

```
    chars.append(char)
```

```
    for ii in range(size):
```

```
        char, h = predict(model, chars[-1], h, top_k=top_k)
```

```
        chars.append(char)
```

```
    return ''.join(chars)
```

← Pass the previous character
and get a new one.

```
sample(model, 2000, top_k=5, prime="And Lven said")
```