

Tuning Deep learning Models

- optimizer Hyperparameters

learning rate, batch size, epochs


- Model Hyperparameters

number of layers, hidden units, model-specific

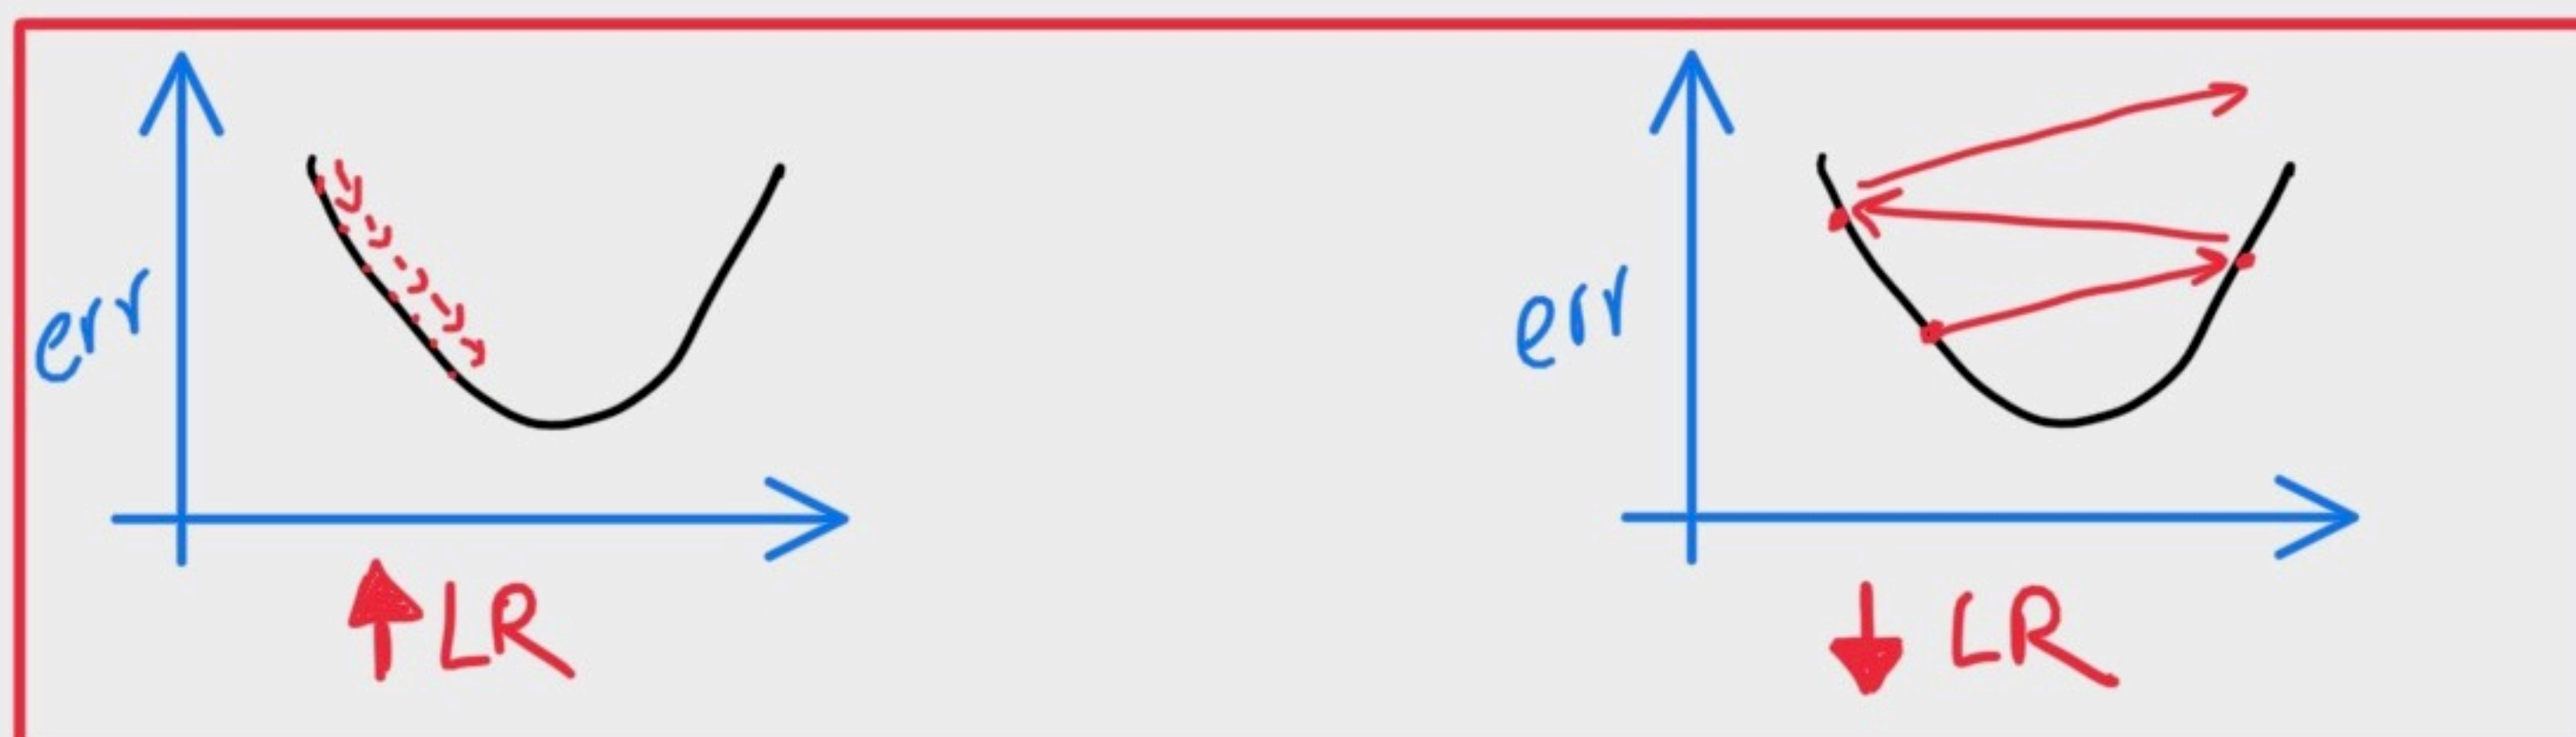
Learning Rate

Seemingly most important hyperparameter.


Good initial value = 0.01



Good Cases
Validation error
decreases



Bad Cases
Validation Error
Increases or
does not decrease
fast enough



Bad Case

Validation Error stops decreasing and somehow oscillates.

Solutions → learning rate decay

- ↓ LR linearly (divide by half every k epochs)
- ↓ LR exponentially (mult by 0.5 every k epochs)
- Adaptive Learning → ↓/↑ as needed


Batch Size

Affects both resource requirements and training

A good initial value $\sim 32, 64$

↑ Batch Size requires more memory.


↓ Batch Size has more noise which helps the GD not to get stuck in local minima.



epochs

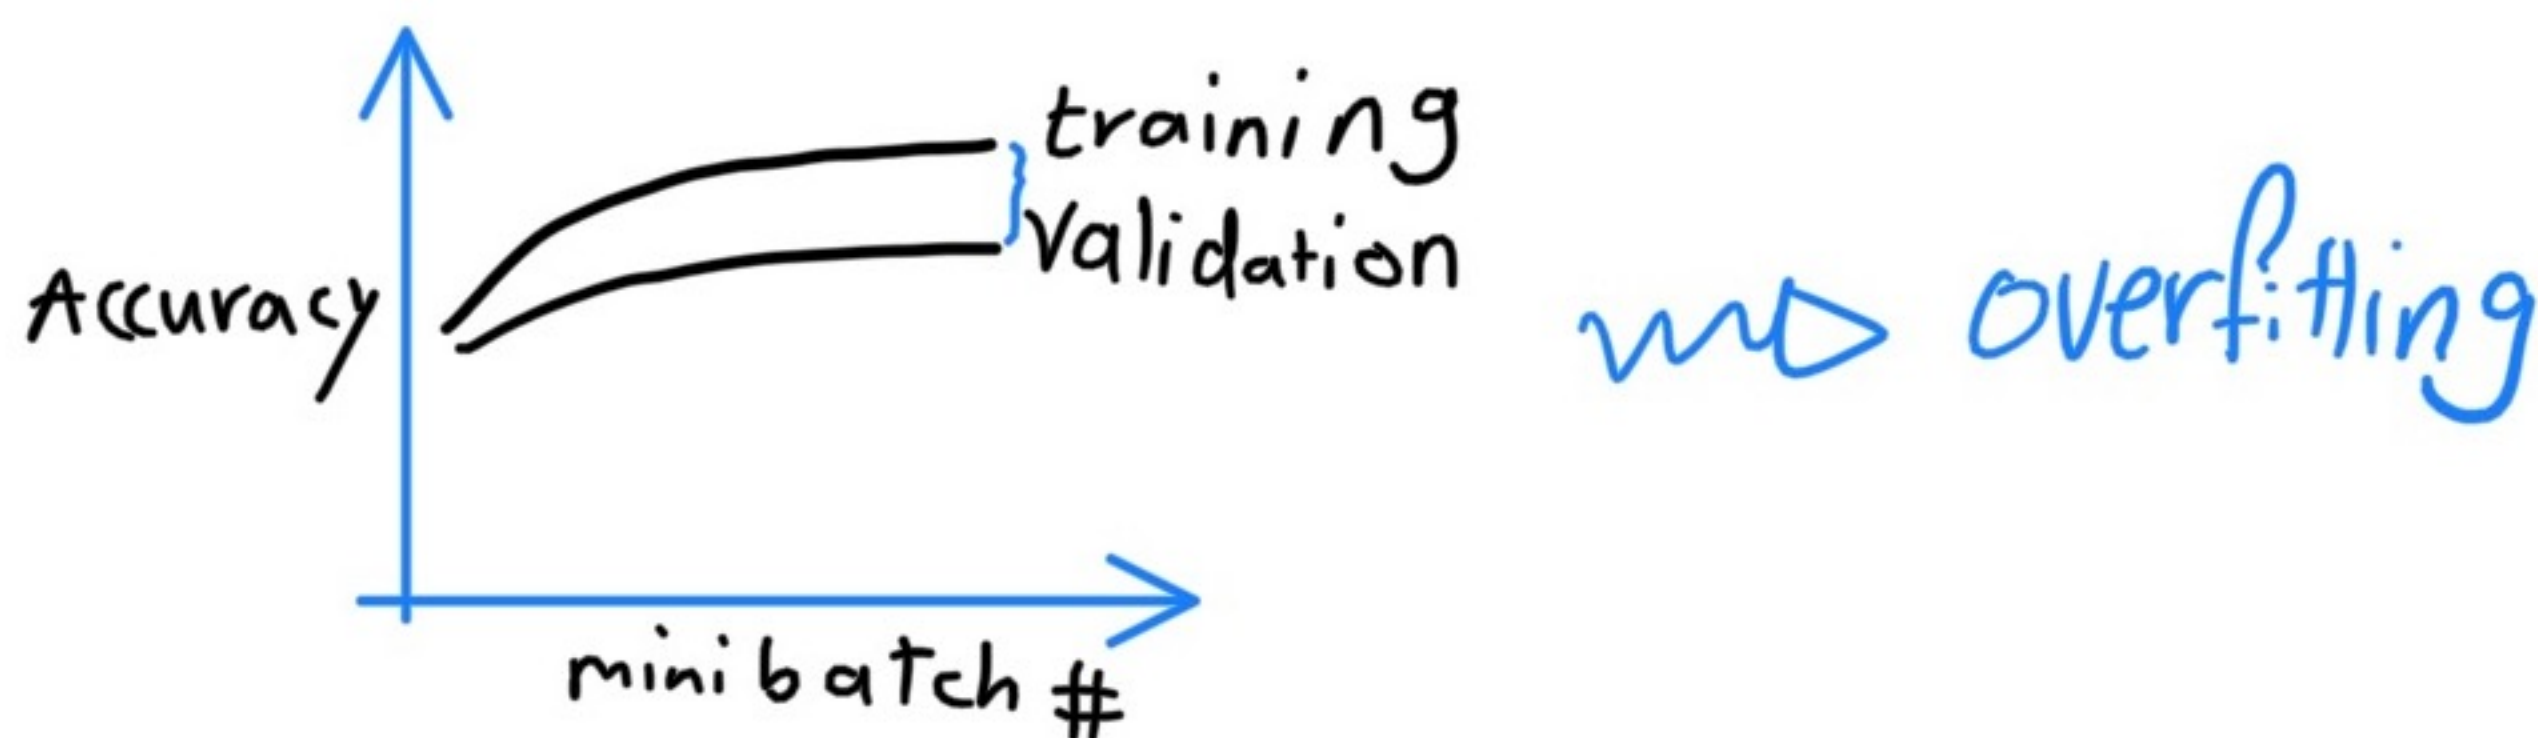
- we must monitor validation loss
when it stops decreasing, we should stop.

- early stopping technique: stop training if the validation loss has not decreased in k epochs.



Model Size

- number of hidden units. Intuitively, it controls the model's capacity to learn a function.
- Having too much capacity \rightarrow overfit



- But the general rule is "More is Better"
- Number of layers. As a rule of thumb, 3 layers is better than 2. But more than that does not usually help. (exception: CNN)

Cell Type

- In an RNN, we can choose LSTM, GRU, or just the vanilla RNN. The first 2 are usually better.
- Choosing LSTM or GRU is task-dependent.

Selecting Reasonable Hyperparameters

training loss \ll validation loss ?

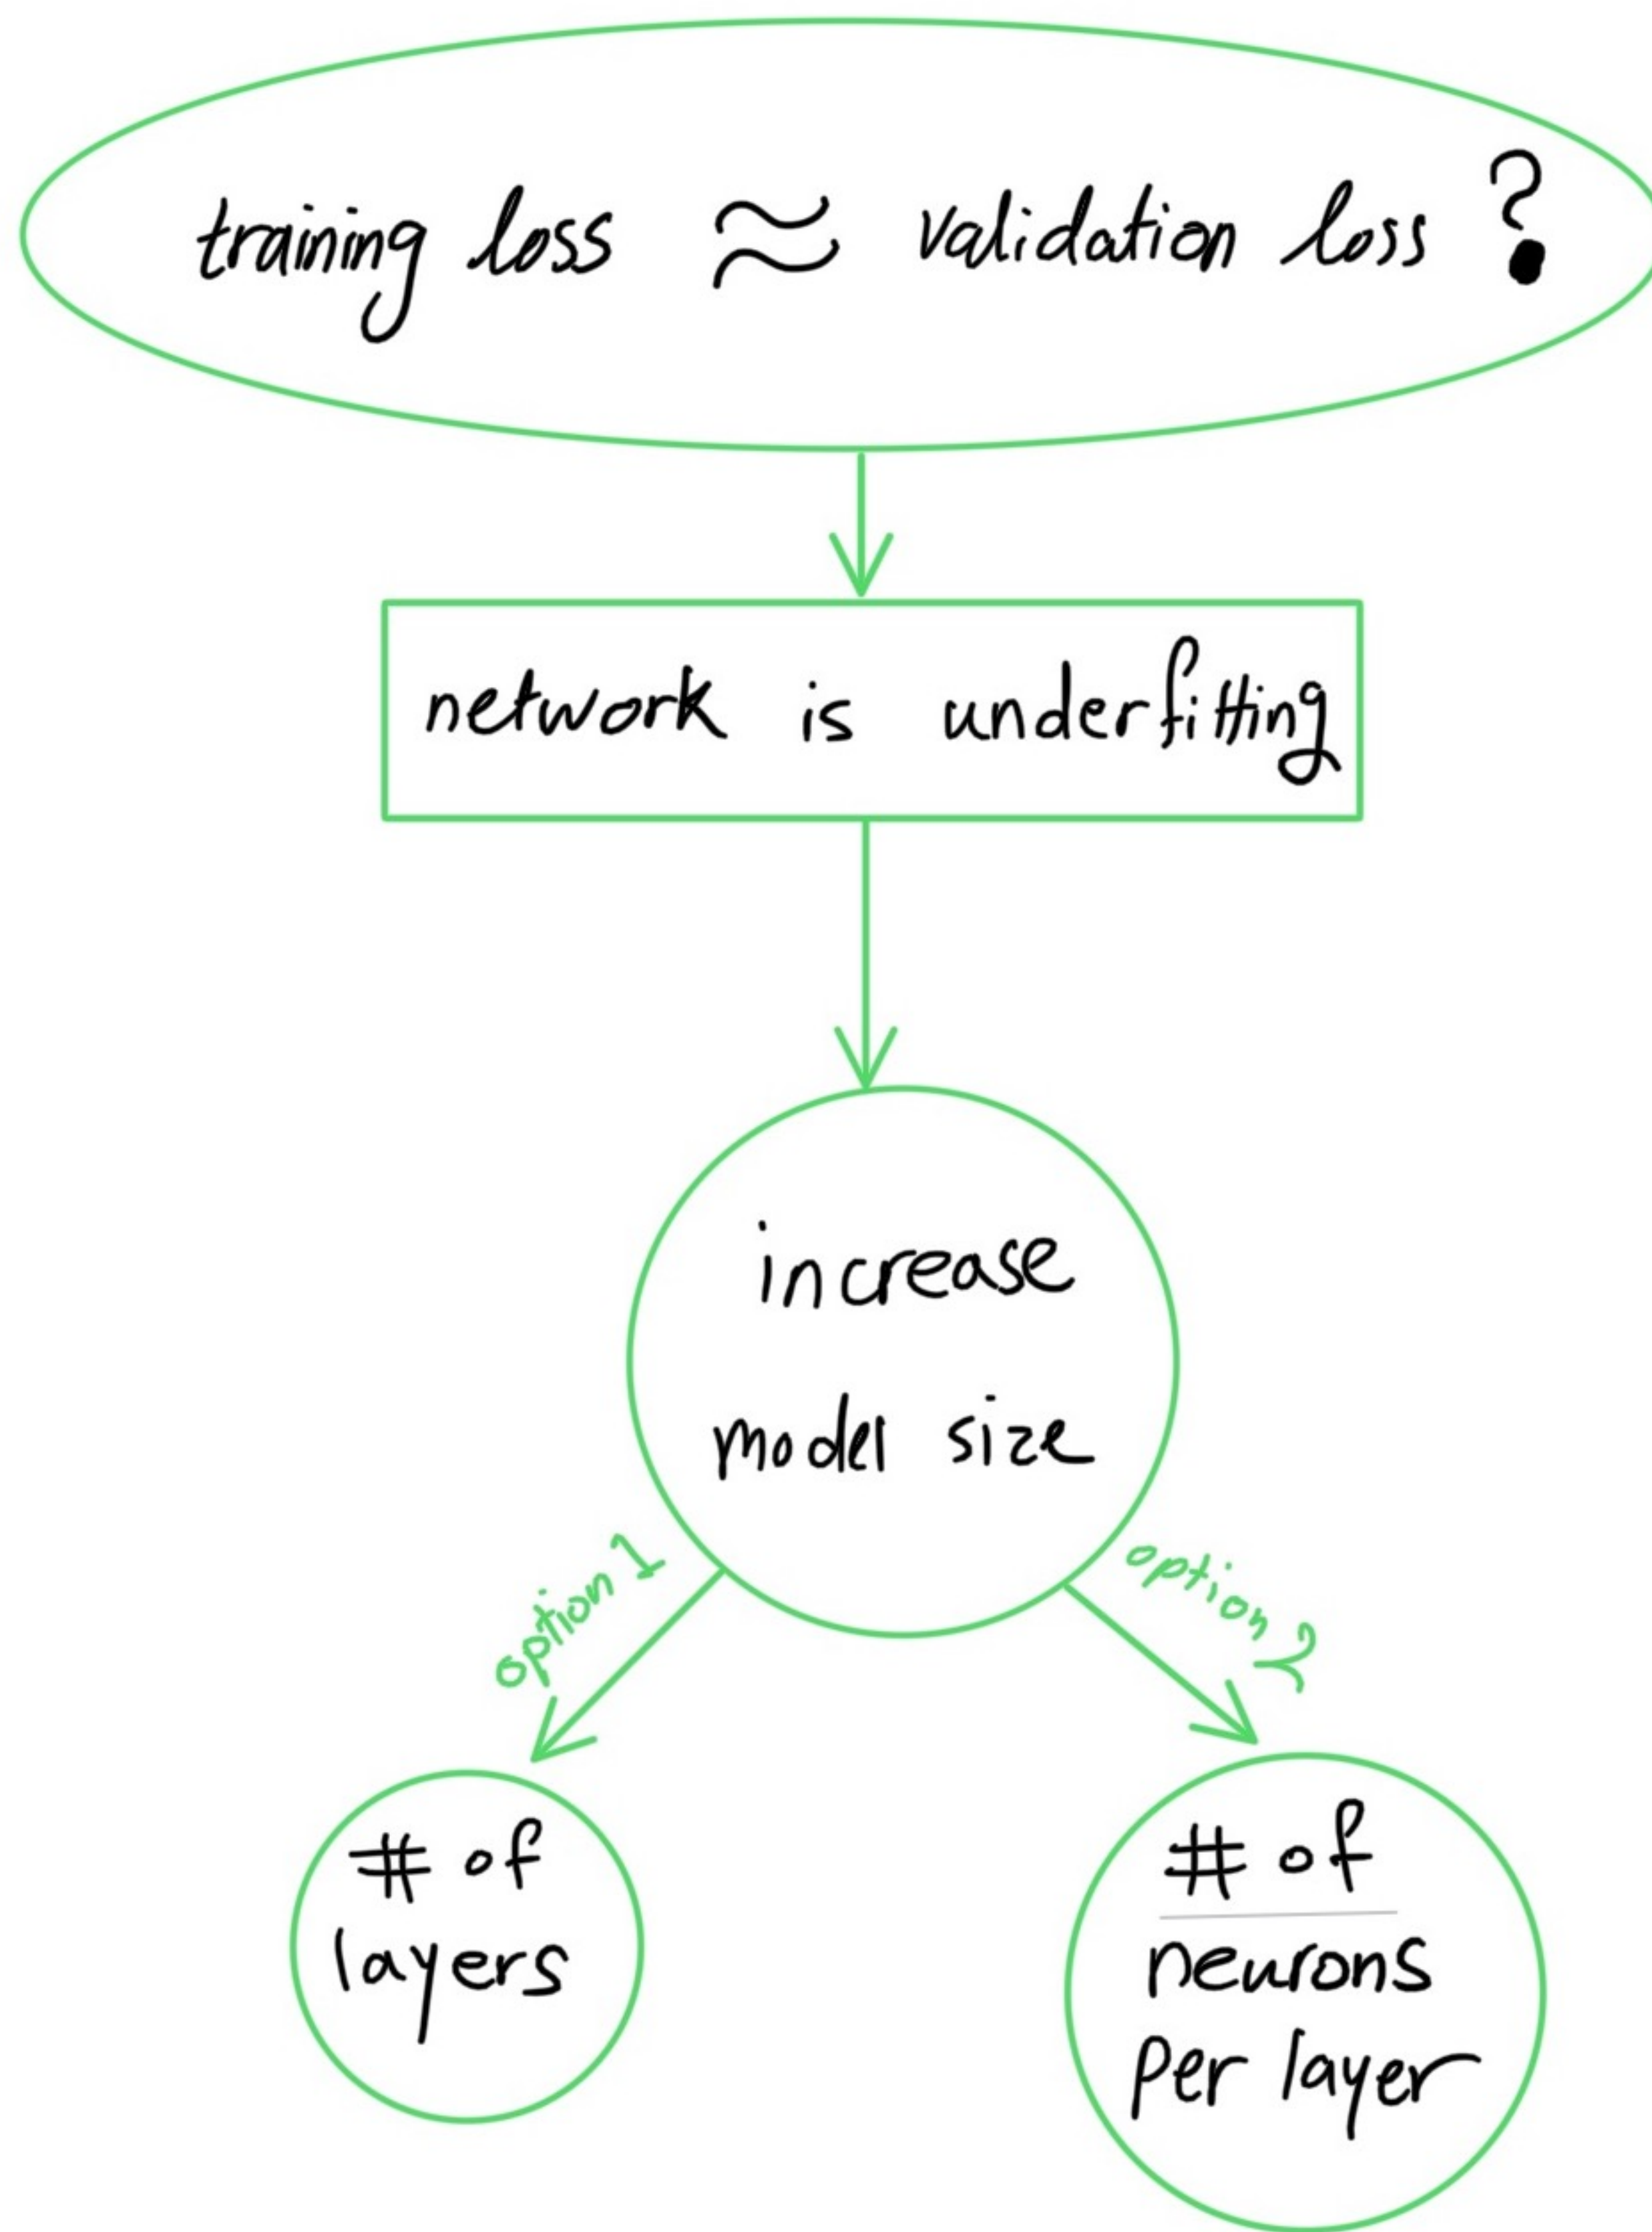
network is overfitting

option 1

increase
drop out

option 2

decrease
model size



- the number of model parameters should be about same magnitude as the size of dataset.

100 MB dataset (\sim 100 million chars)

number of model params \ll size of dataset ?

the model underfits