

Autograd (Backpropagation)

track → $x = \text{torch.zeros}(1, \text{requires_grad=True})$
[with $\text{torch.no_grad}()$:
turns off autograd]

we can get the grad_fn with $x.\text{grad_fn}$
finally, we do $z.\text{backward}()$

Loss and Autograd together:

Code as before
:
 $\text{loss} = \dots$

$\text{loss}.\text{backward}()$ → compute gradients

Using gradients: $\xrightarrow{\text{forward pass}} \text{optimizer}.\text{zero_grad}()$
we use optimizers : $\begin{cases} \text{optimizer}.\text{step}() \\ \hookrightarrow \text{after getting the loss} \end{cases}$

Summary of Training Steps:

1. forward pass
2. compute loss
3. $\text{loss}.\text{backward}()$ to compute gradients
4. take a step with optimizer to update weights

Full training code:

criterion = nn.NLLLoss()
optimizer = optim.SGD(model.parameters(), lr=0.003)

for e in range(epochs):

 running-loss = 0

 for images, labels in trainloader:

 images = images.view(images.shape[0], -1)

 * $\text{optimizer}.\text{zero_grad}()$

 Forward pass

 * $\text{loss}.\text{backward}()$

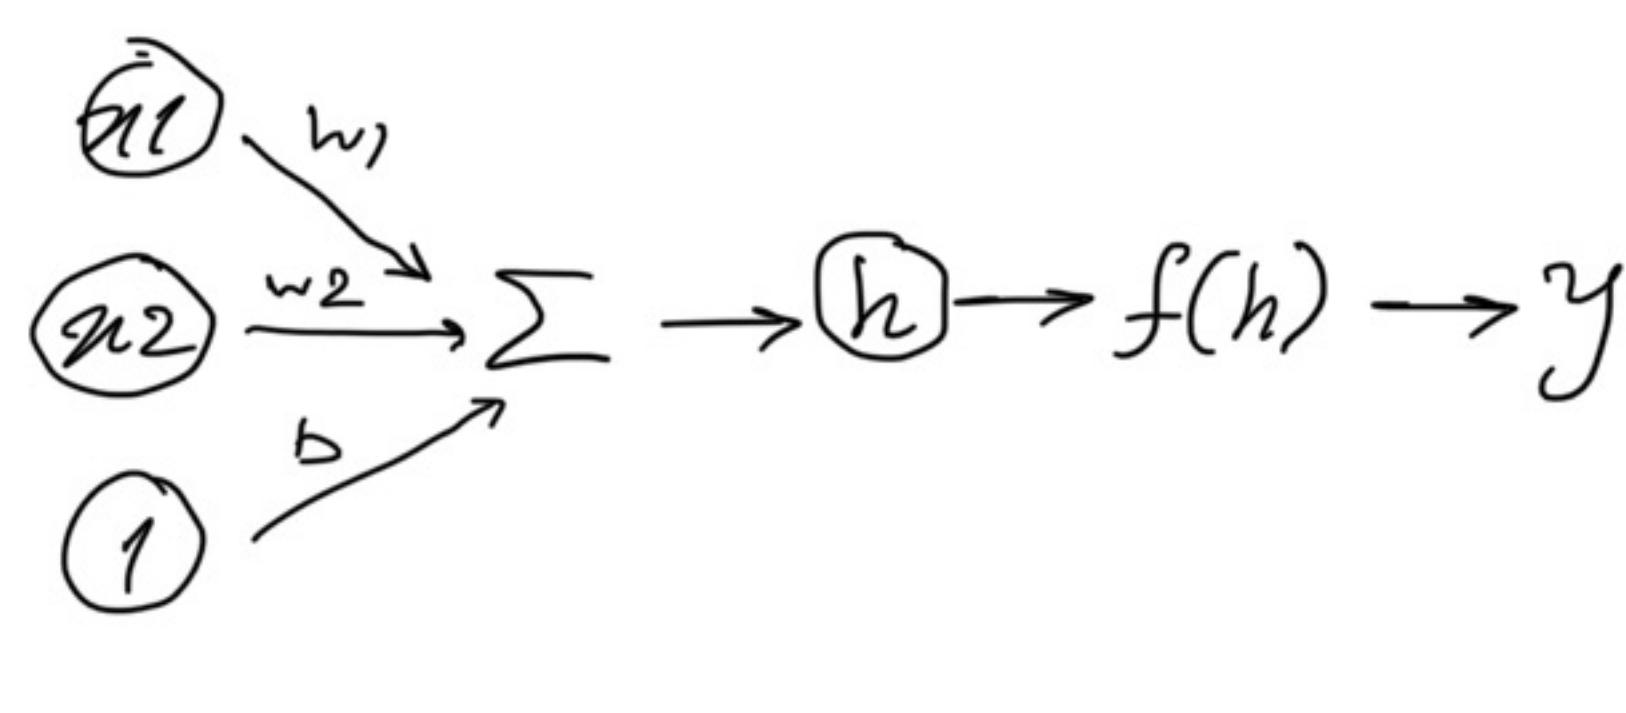
 * $\text{optimizer}.\text{step}()$

 running-loss += loss.item()

 else:

 training-loss = running-loss / len(trainloader)

Tensors



$$y = f(w_1 x_1 + w_2 x_2 + b)$$

$$h = [x_1, \dots, x_n] \cdot \begin{bmatrix} w_1 \\ \vdots \\ w_n \end{bmatrix}$$

def sigmoid(x): return 1/(1+torch.exp(-x))

torch.manual_seed(7)

features = torch.randn((1, 5)) \leftarrow five random normal vars

weights = torch.randn_like(features) \leftarrow a vector with
the shape of features, random.

bias = torch.randn((1, 1)) \leftarrow a 1x1 bias term

elem-by-elem mult

$y = \text{sigmoid}(\text{torch.sum}(\text{features} * \text{weights}) + \text{bias})$

output of a single layer

torch.mm and torch.matmul \rightarrow mat multiplication

for changing the shape of tensors:

1. weights.reshape(a, b) \rightarrow returns new tensor

2. weights.resize_(a, b) \rightarrow in place operation

3. weights.view(a, b) \rightarrow return new data with same data

alt

$y = \text{sigmoid}(\text{torch.mm}(\text{features}, \text{weights.view}(5, 1)) + \text{bias})$

Stacking up units:

$$\vec{h} = [h_1, h_2] = [x_1, \dots, x_n] \cdot \begin{bmatrix} \overset{\text{input units}}{\downarrow} & \overset{\text{hidden units}}{\rightarrow} \\ w_{11} & w_{12} \\ \vdots & \vdots \\ w_{n1} & w_{n2} \end{bmatrix}$$

$h = \text{sigmoid}(\text{torch.mm}(\text{features}, w_1) + b_1)$

$y = \text{output} = \text{sigmoid}(\text{torch.mm}(h, w_2) + b_2)$

Going between Torch and Numpy.

memory is shared between them

$a = np.random.rand(4, 3)$

$b = \text{torch.from_numpy}(a)$

$b.numpy()$

Inference and Validation

we had this code for testing the prediction:

```
model = Classifier()
images, labels = next(iter(testloader))
ps = torch.exp(model(images))
top_p, top_class = ps.topk(1, dim=1)
equals = top_class == labels.view(*top_class.shape)
```

accuracy = torch.mean(equals.type(torch.FloatTensor))

because
equals is a bit tensor
and we have to convert it.

Implementing Validation Pass:

```
for e in epochs:
```

```
    for image, label in trainloader:
```

as before

```
else:
```

the
Validation
pass

```
    with torch.no_grad():
```

above
code

Using nn module to build neural nets :

```
from torch import nn
```

```
class Network(nn.Module):  
    def __init__(self):  
        super().__init__()  
        self.hidden = nn.Linear(784, 256)  
        self.output = nn.Linear(256, 10)
```

- * self.sigmoid = nn.Sigmoid
- * self.softmax = nn.Softmax(dim=1)

```
def forward(self, x):  
    x = self.hidden(x)  
    x = self.sigmoid(x)  
    x = self.output(x)  
    x = self.softmax(x)  
    return x
```

For making things easier: import torch.nn.functional as F
then we can remove self.Sigmoid and self.Softmax
from the class above.

```
def forward(self, x):  
    x = F.sigmoid(self.hidden(x))  
    x = F.softmax(self.output(x), dim=1)  
    return x
```

Avoiding Overfitting:

« dropout »

```
class Classifier(nn.Module):
```

```
    def __init__(self):
```

====

0.2

drop probability ← self.dropout = nn.Dropout($p=0.2$)

```
    def forward(self, x):
```

=====

apply
dropout ← {
at every layer
hidden } $x = \text{self.dropout}(\text{F.relu}(\text{self.fc1}(x)))$
 $x = \text{self.dropout}(\dots)$

no dropout ↴ $x = \text{F.log-softmax}(\text{self.fc4}(x), \text{dim}=1)$

for output layer

return x

Note: when we do validation, we don't want any dropouts, so:

with torch.no_grad():

model.eval() → turns off dropouts

=====

model.train() → revert to train mode

Saving and Loading Models:

all weights and biases of model

Save [`torch.save(model.state_dict(), 'checkpoint')`]

Load [`state_dict = torch.load('checkpoint')`
`model.load_state_dict(state_dict)`]

When loading a checkpoint, it also needs other info such as `input_size`, `output_size`, and number of hidden-layers in addition to the `state_dict`. So:

Saving checkpoint { `checkpoint = { 'input_size': 784,`
 `:`
 `'state_dict': model.state_dict() }`
`torch.save(checkpoint, 'checkpoint')` }

loading checkpoint { `def load_checkpoint(filepath):`
 `checkpoint = torch.load(filepath)`
 `model = fc_model.Network(checkpoint['input'],`
 `- - - >`
 `model.load_state_dict(checkpoint['state_dict'])`
 `return model` }

Loading Image Data:

```
dataset = datasets.ImageFolder('Path',  
                                transform=transforms)
```

image folder structure :

root/dog/m.png

one folder

root/dog/m~.png

for each

:

root/cat/m.png

class

root/cat/m~.png

:

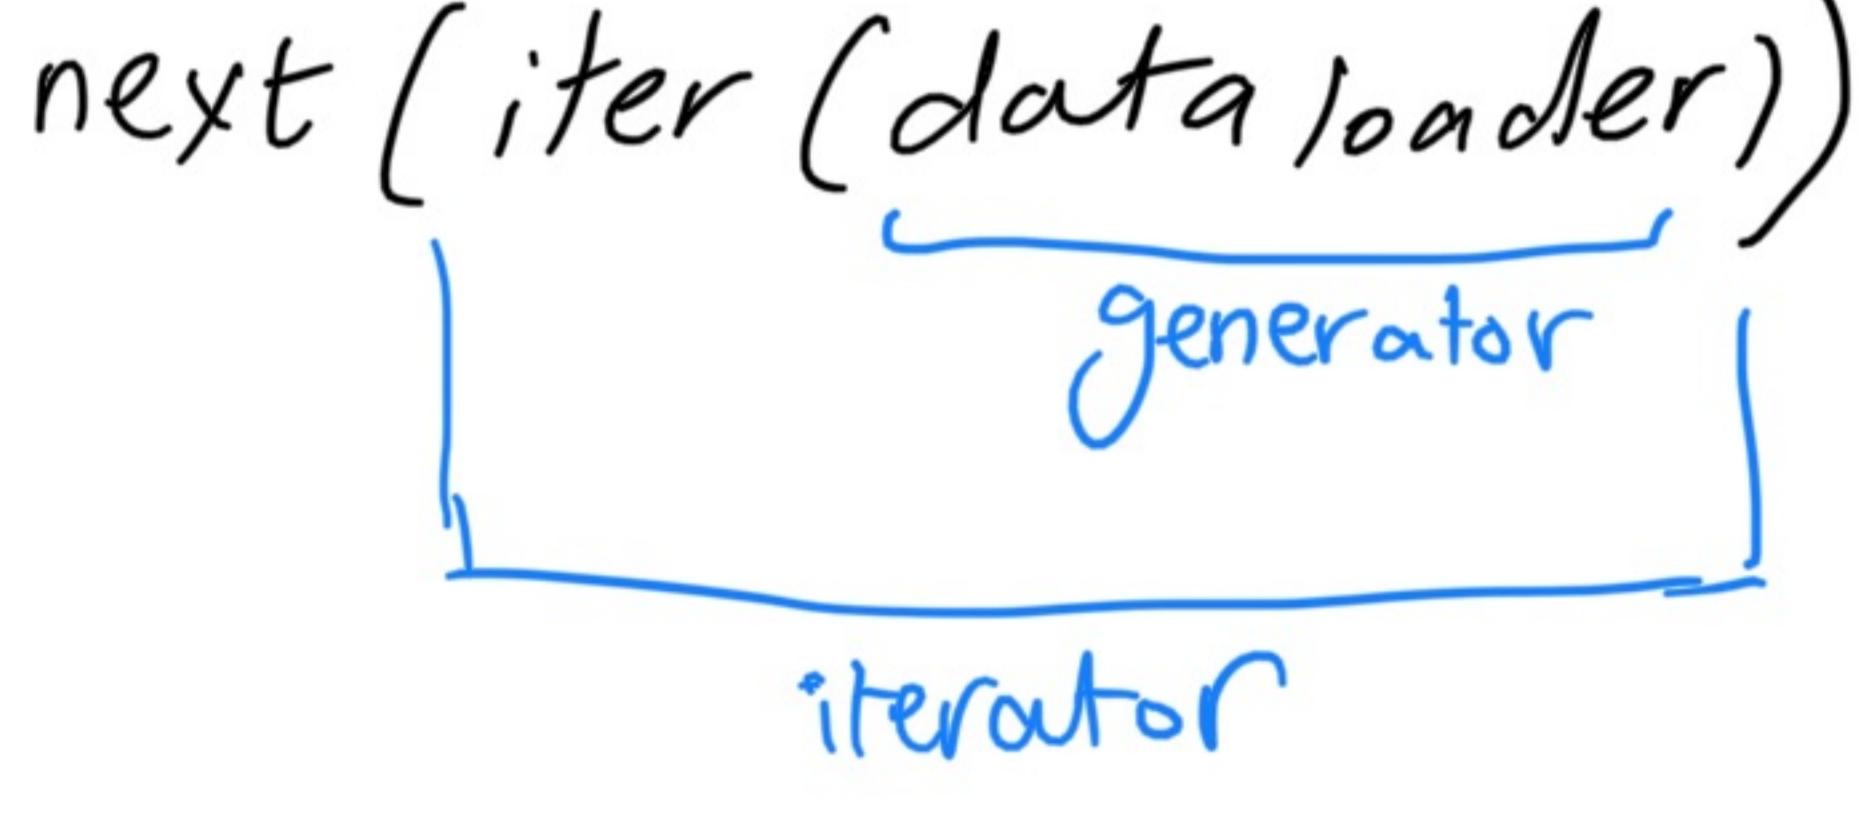
Transforms :

```
transforms = transforms.Compose([  
    transforms.Resize(255),  
    transforms.CenterCrop(224),  
    transforms.ToTensor()])
```

Converting to tensor ← transforms.ToTensor()

Data Loaders :

```
data_loader = torch.utils.data.DataLoader  
(dataset, batch_size=32,  
shuffle=True)
```



Data Augmentation :

Introducing randomness in input data itself.

transforms ← RandomRotation, RandomResizedCrop, etc.

Normalization :

transforms.Normalize

$\text{input[channel]} = (\text{input[channel]} - \text{mean[channel]}) / \text{std[channel]}$

Helps keep data near zero, making backpropagation more stable.

Transfer Learning :

The use of pre-trained models.

`torch ← models.densenet121(pretrained = True)`

then, we have to freeze model params :

`for param in model.parameters():` } we don't
 `param.requires_grad = False` want to
 affect
 these.

Replace their classifier with our classifier:

```
from collections import OrderedDict
classifier = nn.Sequential(OrderedDict([
    ('fc1', nn.Linear(1024, 500)),
    ('relu', nn.ReLU()),
    ('fc2', nn.Linear(500, 2)),
    ('output', nn.LogSoftmax(dim=1))
]))
```

a new,
untrained
classifier

`model.classifier = classifier`

Using GPU :

`model.cuda()` # move the model to GPU
`images.cuda()` # move tensor to gpu

`model.cpu()` } returning the computations
`images.cpu()` to cpu.

Putting Everything Together:

```
device = torch.device("cuda"  
                     if torch.cuda.is_available()  
                     else  
                     "cpu")
```

```
model = models.resnet50(pretrained=True)
```

```
freeze } for param in model.parameters():  
         param.requires_grad = False
```

replace
the pre-
trained
classifier }

```
classifier = nn.Sequential(nn.Linear(2048, 512),  
                           nn.ReLU(),  
                           nn.Dropout(p=0.2),  
                           nn.Linear(512, 2),  
                           nn.LogSoftmax(dim=1))
```

```
model.fc = classifier
```

```
criterion = nn.NLLLoss()
```

```
optimizer = optim.Adam(model.fc.parameters(),  
                       lr=0.003)
```

move the
model to ←
the available

device.

set
hyper
params }

```
epochs = 1  
steps = 0  
running_loss = 0  
print_every = 5
```

Start → for epoch in range(epochs):
actual
training for ims, labs in trainloader:
 steps += 1
 move to ← ims, labs = ims.to(device), labs.to(device)
 the available device

```
optimizer.zero_grad()
```

```
logps = model(images)
```

```
loss = criterion(logps, labels)
```

```
loss.backward()
```

```
optimizer.step()
```

```
running_loss += loss.item()
```

frequency → if step % print_every == 0:
for validation model.eval()
 test_loss = 0
 accuracy = 0

remember to transfer
to gpu:
ims, labs = ims.to(device)
 labs.to(device)

```
for ims, labs in testloader:  
    logps = model(ims)  
    loss = criterion(logps, labs)  
    test_loss += loss.item()
```

ps = torch.exp(logps)
top_ps, top_class = ps.topk(1,
 dim=1)

equality = top_class ==
 labels.view(*top_cls.shape)

accuracy += torch.mean(
 equality.type(torch.FloatTensor)) . item()

running_loss = 0

model.train()

Training Neural Networks

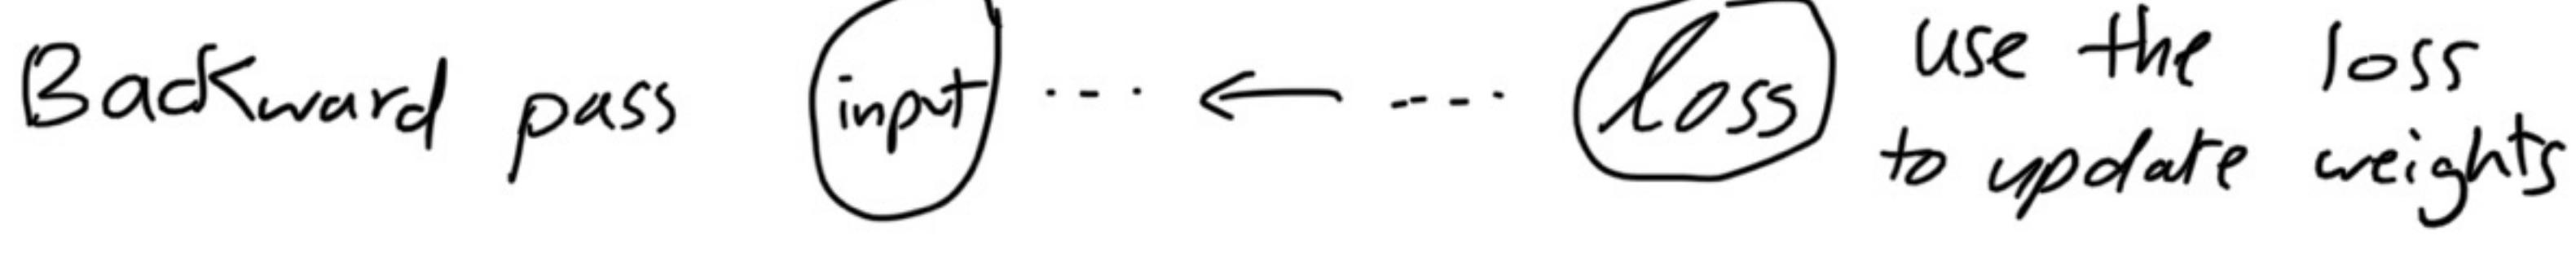
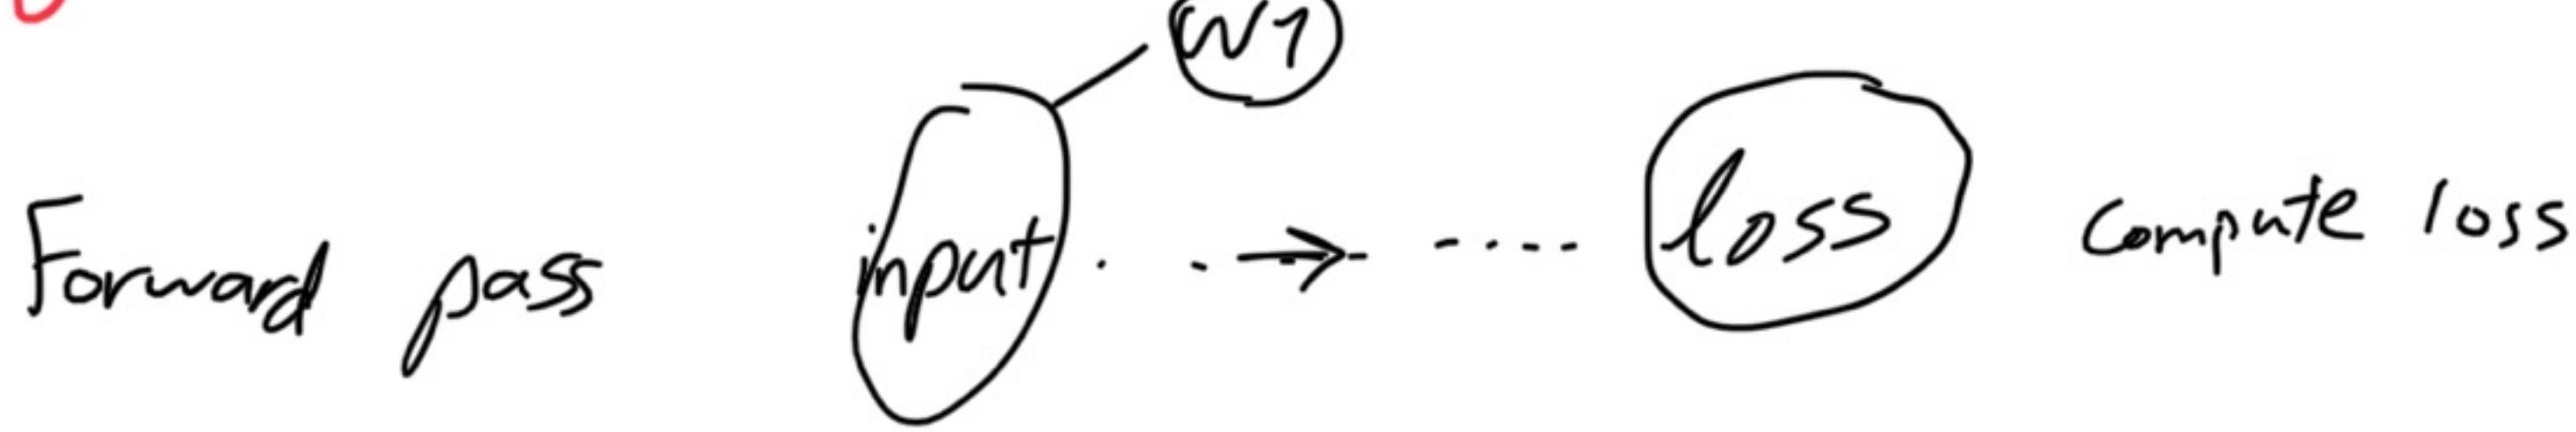
loss function → measure prediction error

$$l = \frac{1}{2n} \sum_i^n (y_i - \hat{y}_i)^2$$

number of training examples

true labels

predicted labels



Loss in Pytorch:

`nn.CrossEntropyLoss` ↗
 ↗ `nn.LogSoftmax()`
 ↘ `nn.NLLLoss()`

input: Scores (not probabilities)

In summary

1. Build a feed forward network:

`model = nn.Sequential(nn.Linear(784, 128),
 nn.ReLU(),
 nn.Linear(128, 64),`

ReLU is almost
always used for
hidden layers

← `nn.ReLU,`

`nn.Linear(64, 10))`

→ we can add a
`nn.LogSoftmax(dim=1)`

2. Define the loss :

and set this
to `nn.NLLLoss()`

`criterion = nn.CrossEntropyLoss()`

3. Get the data :

`images, labels = next(iter(trainloader))`

4. flatten images :

`images = images.view(images.shape[0], -1)`

5. forward pass — get logits

`logits = model(images)`

6. Calculate the loss :

`loss = criterion(logits, labels)`

Neural Networks with Pytorch

```
from torchvision import datasets, transforms
```

for
normalizing transform = transforms.Compose([transforms.ToTensor(),
data
transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

download

mnist
dataset

```
trainset = datasets.MNIST('MNIST-data/',  
download=True,  
train=True,  
transform=transform)
```

define a

loader trainloader = torch.utils.data.DataLoader(trainset,
each time we get a batch, it's shape is (64, 1, 28, 28)
batch-size = 64, shuffle = True)
64 images → 1 color channel → image size

```
dataiter = iter(trainloader)  
images, labels = dataiter.next()
```

we usually flatten the input:

from (64, 1, 28, 28) to (64, 784)

Build a network for identifying images:

inputs = images.view(images.shape[0], -1)

↑
input units

w1 = torch.randn(784, 256)

b1 = torch.randn(256) → hidden units

w2 = torch.randn(256, 10)

b2 = torch.randn(10) → one output digit for each digit.

one hidden layer

h = sigmoid(torch.mm(inputs, w1) + b1)

out = torch.mm(h, w2) + b2

choose the appropriate size

```
def softmax(x):  
    return torch.exp(x) / torch.sum(torch.exp(x), dim=1)  
    • view(-1, 1)
```

probs = softmax(out) # shape = (64, 10)